

Tutorial 05: Parallel Processing - Concurrent Agent Execution

Difficulty: intermediate

Reading Time: 1 hour


Tags: intermediate, parallel, concurrent, performance, multi-agent

Description: Execute multiple agents simultaneously using `ParallelAgent` for faster processing and independent task execution.

Tutorial 05: Parallel Processing - Run Multiple Agents Simultaneously

Overview

Learn how to execute multiple agents concurrently to dramatically speed up your workflows! This tutorial teaches you the `ParallelAgent` pattern and the powerful **fan-out/gather** technique - essential for any real-world agent system that needs performance.

 **Working Implementation Available:** A complete, tested travel planning system is available at `tutorial_implementation/tutorial05/` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/). The implementation includes comprehensive tests, documentation, and a user-friendly setup process.

Prerequisites

- **Completed Tutorials 01-04** - Understanding of agents, tools, and sequential workflows
- **Installed ADK** - `pip install google-adk`
- **API key configured** - From Tutorial 01

Core Concepts

| ParallelAgent

The `ParallelAgent` executes multiple sub-agents **concurrently** (at the same time), not sequentially. This is perfect when:

- Tasks are **independent** (don't depend on each other)
- You want **speed** (faster than running one-by-one)
- You're gathering information from **multiple sources**

Key Characteristics:

- All sub-agents start at the same time
- Agents run independently (no shared data during execution)
- `ParallelAgent` waits for ALL to complete
- Returns when the slowest agent finishes

| When to Use Parallel vs Sequential

Use `ParallelAgent` when:

- ✓ Tasks are independent (can run in any order)
- ✓ Speed matters (want results faster)
- ✓ Gathering data from multiple sources
- ✓ Tasks don't need each other's outputs

Use `SequentialAgent` when:

- ✓ Tasks MUST happen in specific order

- ✓ Each step needs previous step's output
- ✓ Building a pipeline (output flows between agents)

Fan-Out/Gather Pattern

This is THE killer pattern for real-world agents:

1. **Fan-Out** (Parallel) → Run multiple agents concurrently to gather data
2. **Gather** (Sequential) → Merge all results with a final agent



ParallelAgent (fast!)

SequentialAgent (combine)

This gives you both **speed** (parallel data gathering) and **synthesis** (merging results).

Use Case

We're building a **Smart Travel Planner** that helps users plan trips by:

1. **Searching flights** (concurrent)
2. **Searching hotels** (concurrent)
3. **Finding activities** (concurrent)
4. **Merging everything** into a complete itinerary (sequential)

This demonstrates the fan-out/gather pattern - parallel data gathering + sequential synthesis!

Step 1: Get the Working Implementation

A complete, tested implementation is available in the repository:

```
# Navigate to the working implementation
cd tutorial_implementation/tutorial05/

# Install dependencies
make setup

# Copy environment template and add your API key
cp travel_planner/.env.example travel_planner/.env
# Edit travel_planner/.env and add your GOOGLE_API_KEY
```

Alternative: Follow the step-by-step build instructions below to create your own implementation.

Step 2: Create Project Structure (Optional - Skip if using working implementation)

If you prefer to build from scratch, create this structure:

```
mkdir travel_planner
cd travel_planner
touch __init__.py agent.py .env
```

Copy your `.env` file from previous tutorials.

Project Structure

The working implementation follows this structure:

```
tutorial05/
├── travel_planner/           # Agent implementation
│   ├── __init__.py         # Package initialization
│   ├── agent.py            # Agent definitions and pipeline
│   └── .env.example         # Environment template
├── tests/                   # Comprehensive test suite
│   ├── __init__.py
│   ├── test_agent.py       # Agent and pipeline tests
│   ├── test_imports.py     # Import validation tests
│   └── test_structure.py   # Project structure tests
├── requirements.txt         # Python dependencies
├── Makefile                # Development commands
└── README.md               # Implementation documentation
```

Key Components:

- `travel_planner/agent.py` : Complete ParallelAgent + SequentialAgent pipeline
- `tests/test_agent.py` : 57 tests covering all functionality
- `Makefile` : `make setup`, `make test`, `make dev`, `make demo` commands
- `.env.example` : Template for API key configuration

Step 2: Set Up Package Import

travel_planner/init.py

```
from . import agent
```

Step 3: Define the Travel Planning Agents

travel_planner/agent.py

```

from __future__ import annotations

from google.adk.agents import Agent, ParallelAgent, SequentialAgent

# =====
# PARALLEL SEARCH AGENTS
# =====

# ===== Parallel Branch 1: Flight Finder =====
flight_finder = Agent(
    name="flight_finder",
    model="gemini-2.0-flash",
    description="Searches for available flights",
    instruction=(
        "You are a flight search specialist. Based on the user's travel request\n"
        "search for available flights.\n"
        "\n"
        "Provide 2-3 flight options with:\n"
        "- Airline name\n"
        "- Departure and arrival times\n"
        "- Price range\n"
        "- Flight duration\n"
        "\n"
        "Format as a bulleted list. Be specific and realistic."
    ),
    output_key="flight_options" # Saves to state
)

# ===== Parallel Branch 2: Hotel Finder =====
hotel_finder = Agent(
    name="hotel_finder",
    model="gemini-2.0-flash",
    description="Searches for available hotels",
    instruction=(
        "You are a hotel search specialist. Based on the user's travel request\n"
        "find suitable hotels.\n"
        "\n"
        "Provide 2-3 hotel options with:\n"
        "- Hotel name and rating\n"
        "- Location (district/area)\n"
        "- Price per night\n"
        "- Key amenities\n"
        "\n"
        "Format as a bulleted list. Be specific and realistic."
    ),
    output_key="hotel_options" # Saves to state
)

```

```

)

# ===== Parallel Branch 3: Activity Finder =====
activity_finder = Agent(
    name="activity_finder",
    model="gemini-2.0-flash",
    description="Finds activities and attractions",
    instruction=(
        "You are a local activities expert. Based on the user's travel request\n"
        "recommend activities and attractions.\n"
        "\n"
        "Provide 4-5 activity suggestions with:\n"
        "- Activity name\n"
        "- Description (1 sentence)\n"
        "- Estimated duration\n"
        "- Estimated cost\n"
        "\n"
        "Format as a bulleted list. Include mix of paid/free activities."
    ),
    output_key="activity_options" # Saves to state
)

# =====
# FAN-OUT: PARALLEL DATA GATHERING
# =====

# Create the ParallelAgent for concurrent search
parallel_search = ParallelAgent(
    name="ParallelSearch",
    sub_agents=[
        flight_finder,
        hotel_finder,
        activity_finder
    ], # All run AT THE SAME TIME!
    description="Searches flights, hotels, and activities concurrently"
)

# =====
# GATHER: SEQUENTIAL RESULT MERGING
# =====

# ===== Gather: Merge Results into Itinerary =====
itinerary_builder = Agent(
    name="itinerary_builder",
    model="gemini-2.0-flash",
    description="Combines all search results into a complete travel itinerary"
    instruction=(

```

```

        "You are a travel planner. Create a complete, well-organized itinerary
        "by combining the search results below.\n"
        "\n"
        "***Available Flights:**\n"
        "{flight_options}\n" # Reads from state!
        "\n"
        "***Available Hotels:**\n"
        "{hotel_options}\n" # Reads from state!
        "\n"
        "***Recommended Activities:**\n"
        "{activity_options}\n" # Reads from state!
        "\n"
        "Create a formatted itinerary that:\n"
        "1. Recommends the BEST option from each category (flights, hotel)\n"
        "2. Organizes activities into a day-by-day plan\n"
        "3. Includes estimated total cost\n"
        "4. Adds helpful travel tips\n"
        "\n"
        "Format beautifully with clear sections and markdown."
    ),
    output_key="final_itinerary"
)

# =====
# COMPLETE FAN-OUT/GATHER PIPELINE
# =====

# Combine parallel search with sequential merge
travel_planning_system = SequentialAgent(
    name="TravelPlanningSystem",
    sub_agents=[
        parallel_search,      # Step 1: Gather data in parallel (FAST!)
        itinerary_builder     # Step 2: Merge results (synthesis)
    ],
    description="Complete travel planning system with parallel search and itin
)

# MUST be named root_agent for ADK discovery
root_agent = travel_planning_system

```

Code Breakdown

Fan-Out/Gather Flow:


```

User: "Plan a trip to Tokyo for 5 days"
↓
ParallelAgent starts (3 agents run AT SAME TIME):
  └─ flight_finder → searches flights → state['flight_options']
  └─ hotel_finder → searches hotels → state['hotel_options']
  └─ activity_finder → finds activities → state['activity_options']
  ↓ (waits for ALL to complete)
ParallelAgent completes
↓
itinerary_builder runs (reads all 3 keys from state)
↓
Final Output: complete travel itinerary

```

Why This Pattern Works:

1. **Parallel search is FAST** - 3 agents run concurrently (not one-by-one)
2. **Independent tasks** - Flight search doesn't need hotel data
3. **Sequential merge** - Itinerary builder needs ALL search results
4. **Best of both worlds** - Speed (parallel) + synthesis (sequential)

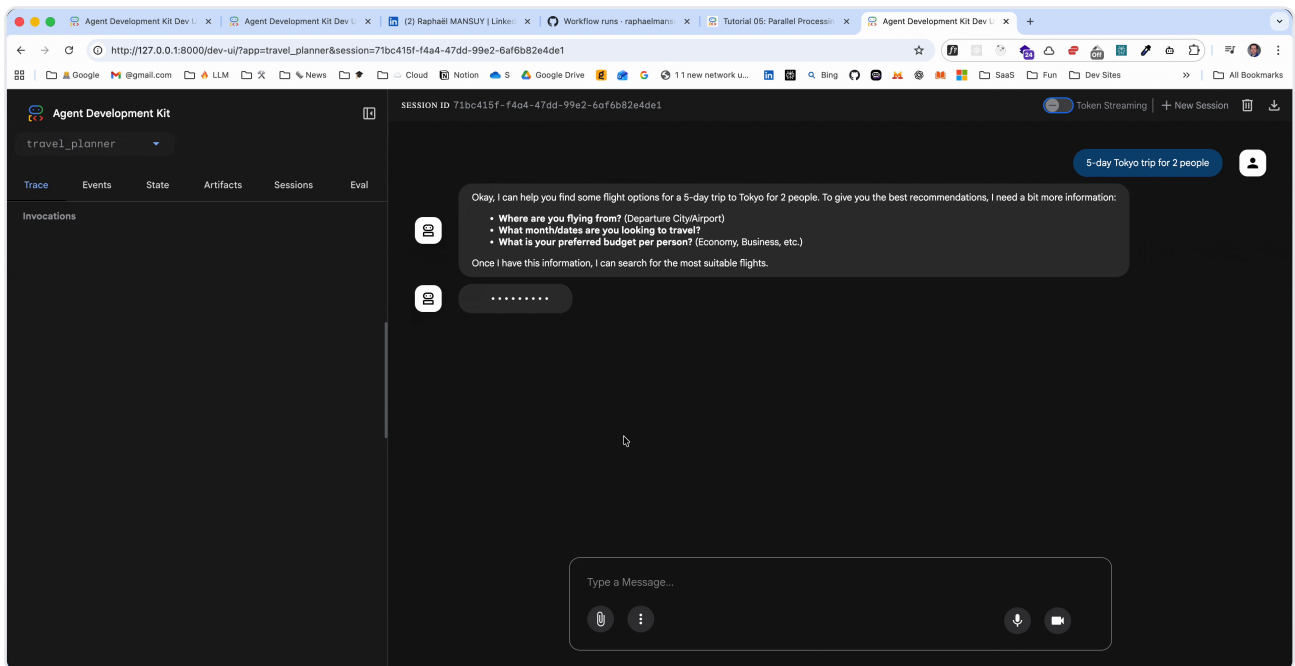
Key Design Decisions:

- Each search agent has focused responsibility
- All search agents save to state with `output_key`
- Itinerary builder reads all state keys with `{key}` syntax
- Outer `SequentialAgent` ensures parallel → sequential flow

Step 4: Run the Travel Planner

| Demo in Action

Here's what your parallel processing travel planner looks like in action:



Using the Working Implementation

```
# From tutorial_implementation/tutorial05/
make dev
```

Open `http://localhost:8000` and select "travel_planner".

Manual Setup (if building from scratch)

Navigate to parent directory and launch:

```
cd .. # Go to parent of travel_planner/
adk web
```

Open `http://localhost:8000` and select "travel_planner".

Try These Prompts

Basic Trip:

```
Plan a 3-day trip to Paris
```

Detailed Request:

I need a 5-day Tokyo trip for 2 people. Budget-friendly options preferred.

Specific Requirements:

Weekend getaway to New York. Looking for cultural activities and good restaurants.

Beach Vacation:

Plan a relaxing week in Bali with beach activities.

Understanding Parallel Execution

Open the **Events tab** and watch the magic:

1. **Event:** ParallelAgent starts
2. **Events (ALL AT ONCE):**
 3. Flight finder starts
 4. Hotel finder starts
 5. Activity finder starts
6. **Events (AS THEY COMPLETE):**
 7. Flight finder completes → saves to state
 8. Hotel finder completes → saves to state
 9. Activity finder completes → saves to state
10. **Event:** ParallelAgent completes (when ALL done)
11. **Event:** Itinerary builder starts (with all 3 keys injected)
12. **Event:** Itinerary builder completes → final output!

Notice: The 3 search agents don't run in order - they all start together and finish whenever they're done!

Testing Your Implementation

The working implementation includes comprehensive tests to validate your understanding:

```
# From tutorial_implementation/tutorial05/  
make test
```

Test Coverage:

- ✓ Agent configurations and instructions (57 tests total)
- ✓ ParallelAgent structure and concurrent execution
- ✓ SequentialAgent pipeline flow and state management
- ✓ Data injection between parallel and sequential agents
- ✓ Import validation and module structure
- ✓ Project organization and file structure

Quick Demo:

```
# Test basic functionality without full ADK setup  
make demo
```

This validates that your agents load correctly and the pipeline structure is sound.

Expected Behavior

Example: "Plan a 3-day trip to Paris"

User: Plan a 3-day trip to Paris

[All 3 agents run in parallel - you'll see them in Events]

Flight Options Found:

- Air France AF1234: NYC → CDG, 7:00am-8:30pm, \$650, 7h 30m
- Delta DL456: NYC → CDG, 10:00am-11:45pm, \$720, 7h 45m

Hotel Options Found:

- Hotel Le Marais (4★): Central Paris, \$180/night, WiFi + Breakfast
- Montmartre Boutique (3★): Artistic district, \$130/night, Rooftop view

Activity Options Found:

- Eiffel Tower: Iconic landmark, 2-3 hours, €25
- Louvre Museum: World-famous art, 3-4 hours, €17
- Seine River Cruise: Scenic evening cruise, 1 hour, €15
- Montmartre Walk: Historic neighborhood, 2 hours, Free
- French Cooking Class: Learn local cuisine, 3 hours, €85

[Itinerary builder merges everything]

Final Itinerary:

Your 3-Day Paris Adventure

Recommended Flight

- ✈️ ****Air France AF1234**** - Best balance of price and timing
- Departs 7:00am, arrives 8:30pm local time
- Price: \$650 per person

Accommodation

- 🏨 ****Hotel Le Marais (4★)**** - Perfect central location
- \$180/night × 3 nights = \$540
- Includes breakfast and WiFi

Day-by-Day Plan

...

How It Works (Behind the Scenes)

Parallel Execution:

1. `ParallelAgent.sub_agents` = [flight, hotel, activity]
2. ADK calls `.run()` on ALL agents simultaneously (async)

3. Each agent executes independently (no shared state during run)
4. ADK waits for ALL to complete
5. Each agent's output saved to its `output_key` in state

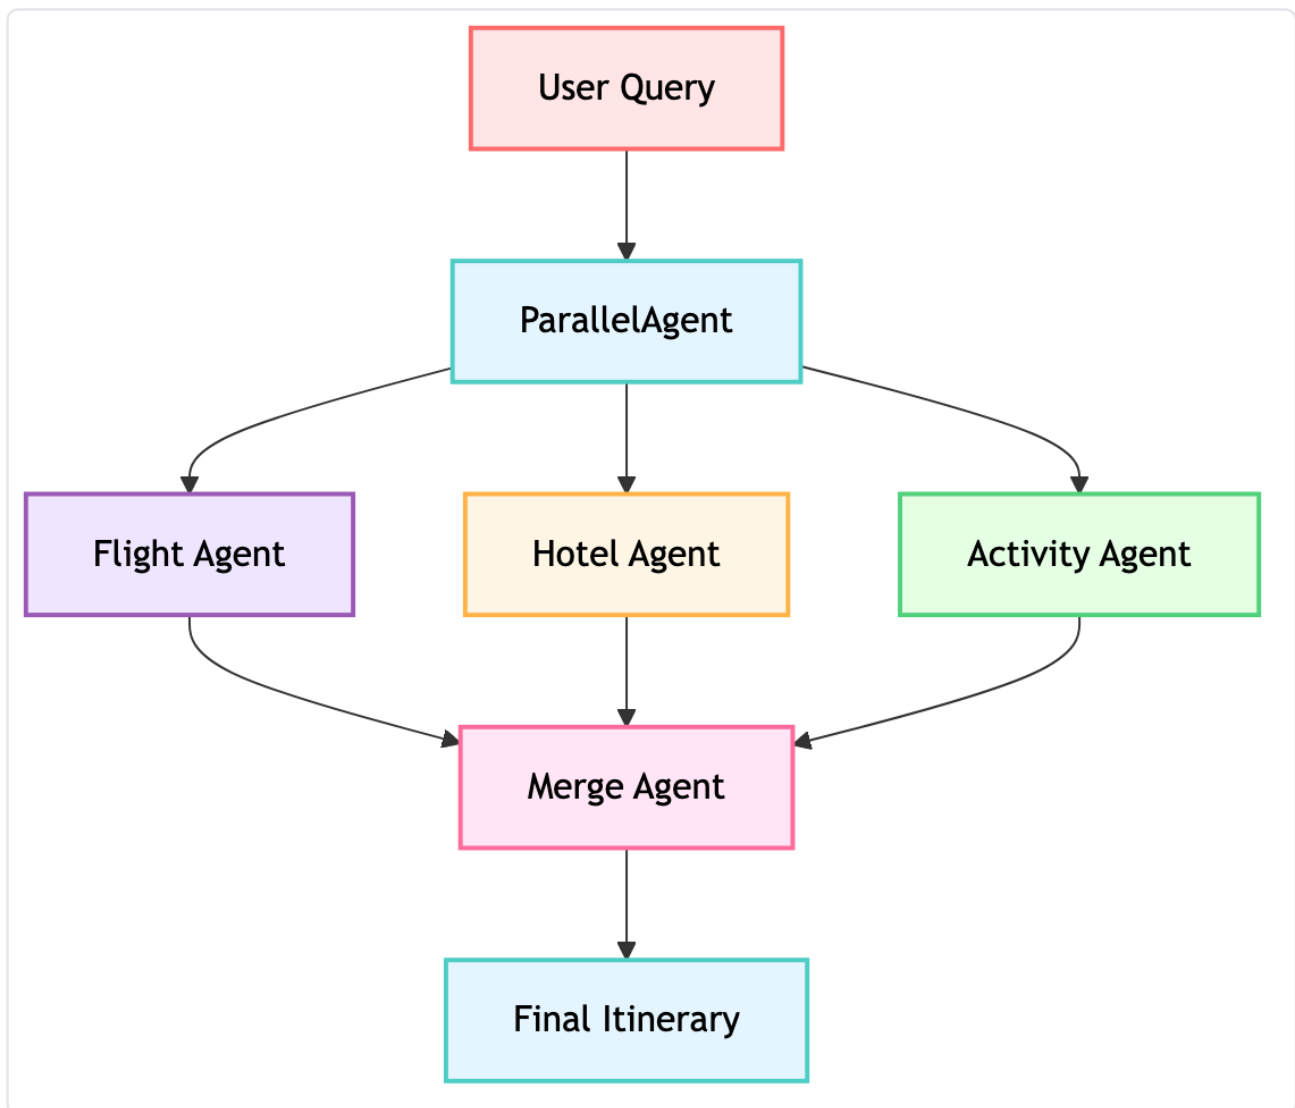
Sequential Merge:

1. After parallel completes, next agent in sequence runs
2. Itinerary builder's instruction has `{flight_options}`, `{hotel_options}`, `{activity_options}`
3. ADK injects values from state into instruction
4. Builder creates final output using all gathered data

Performance:

- Sequential: ~30 seconds (3 agents × 10 seconds each)
- Parallel: ~10 seconds (all 3 run together, limited by slowest)
- **3x faster!**

Visual Flow:



Key Takeaways

- ✓ **ParallelAgent runs sub-agents concurrently** - Major speed boost
- ✓ **Perfect for independent tasks** - Flight search doesn't need hotel data
- ✓ **All agents must complete** - ParallelAgent waits for slowest agent
- ✓ **Fan-out/gather is the pattern** - Parallel for speed, Sequential for synthesis
- ✓ **Combine with SequentialAgent** - Best of both worlds
- ✓ **Watch Events tab** - See agents starting simultaneously
- ✓ **Real performance gains** - 3-10x faster than sequential for I/O-bound tasks

Best Practices

DO:

- Use for I/O-bound tasks (API calls, web searches, database queries)
- Keep sub-agents independent (no dependencies between them)
- Combine with Sequential for merge/synthesis step
- Use descriptive output_keys for each parallel agent
- Monitor Events tab to verify parallel execution

DON'T:

- Use when agents need each other's outputs (use Sequential instead)
- Assume specific execution order within parallel block
- Forget the merge step (parallel gives you data, need synthesis)
- Use for CPU-bound tasks without threading consideration
- Expect perfect linear speedup (limited by slowest agent)

Common Issues

Problem: "Parallel agents seem to run sequentially"

- **Solution:** Check Events tab - they should show same start times
- **Solution:** Might be model API rate limiting

Problem: "Itinerary builder missing some data"

- **Solution:** Verify each search agent has `output_key` defined
- **Solution:** Check `{key}` names match exactly in itinerary instruction

Problem: "One agent fails, whole system stops"

- **Solution:** ParallelAgent waits for ALL - one failure blocks
- **Solution:** Consider error handling in individual agents

Problem: "Not seeing speed improvements"

- **Solution:** Speed gain depends on task type (I/O vs CPU)
- **Solution:** Check if you're being rate-limited by APIs

What We Built

You now have a production-quality travel planner that:

- Searches multiple sources concurrently (3x faster!)
- Gathers comprehensive travel information
- Synthesizes everything into a complete itinerary
- Demonstrates the fan-out/gather pattern


And you understand when/how to use parallel execution!

Real-World Applications

Parallel Execution Is Perfect For:

- **Data Gathering:** Search multiple APIs/databases simultaneously
- **Content Creation:** Generate multiple variations concurrently
- **Analysis:** Run different analyses on same data in parallel
- **Validation:** Check multiple conditions simultaneously
- **Multi-Source Research:** Academic, news, social media, patents
- **Comparison Shopping:** Prices, reviews, availability from different sources

Next Steps

 **Tutorial 06: Multi-Agent Systems** - Combine Sequential and Parallel patterns for sophisticated workflows

Further Reading:

- [Parallel Agents Documentation](https://google.github.io/adk-docs/agents/workflow-agents/parallel-agents/) (https://google.github.io/adk-docs/agents/workflow-agents/parallel-agents/)
- [Workflow Agents Overview](https://google.github.io/adk-docs/agents/workflow-agents/) (https://google.github.io/adk-docs/agents/workflow-agents/)
- [Performance Optimization Guide](https://google.github.io/adk-docs/performance/) (https://google.github.io/adk-docs/performance/)

Exercises (Try On Your Own!)

1. **Add more parallel branches** - Include weather, currency exchange, travel advisories
2. **Error handling** - Handle case where one search fails
3. **User preferences** - Add a preferences agent that customizes recommendations
4. **Nested parallel** - Have each search agent spawn its own parallel sub-searches
5. **Benchmarking** - Measure actual speedup comparing sequential vs parallel

Complete Code Reference

Working Implementation: See `tutorial_implementation/tutorial05/` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/) for a complete, tested version with comprehensive documentation.

Key Files:

- `travel_planner/agent.py` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/travel_planner/agent.py) - Complete agent implementation
- `tests/test_agent.py` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/tests/test_agent.py) - 57 comprehensive tests
- `README.md` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/README.md) - Detailed implementation guide
- `Makefile` (https://github.com/raphaelmansuy/adk_training/tree/main/tutorial_implementation/tutorial05/Makefile) - Development commands

Quick Start with Working Code:

```
cd tutorial_implementation/tutorial05/
make setup # Install dependencies
make test  # Run all tests (57 passing)
make dev   # Start development server
```

Manual Implementation:

`travel_planner/init.py`

```
from . import agent
```

travel_planner/.env

```
GOOGLE_GENAI_USE_VERTEXAI=FALSE  
GOOGLE_API_KEY=your-api-key-here
```

travel_planner/agent.py

```
# See Step 3 above for complete code
```

Congratulations! You've mastered parallel processing and the fan-out/gather pattern!



Generated on 2025-10-19 17:56:17 from 05_parallel_processing.md

Source: Google ADK Training Hub