

# Tutorial 17: Agent-to-Agent Communication - Distributed Systems

---

**Difficulty:** advanced

**Reading Time:** 2 hours

**Tags:** advanced, agent-communication, distributed, delegation, coordination

**Description:** Build distributed agent systems with agent-to-agent communication, delegation, and coordination for complex multi-agent orchestration.

:::info FULLY WORKING A2A IMPLEMENTATION - TESTED & VERIFIED

**This tutorial features a complete, tested A2A implementation using the official Google ADK.**

- ✓ **All Issues Resolved:** Agent card URLs fixed, context handling improved
- ✓ **Working End-to-End:** Complete orchestration with meaningful responses
- ✓ **Tested Implementation:** Real working code with successful test results

**Key approach:** `uvicorn + to_a2a()` for servers, `RemoteA2aAgent` for clients, with proper A2A context handling for intelligent remote agent responses.

**Latest Update:** January 10, 2025 - Fixed context handling for production-ready A2A.

:::

## Tutorial 17: Agent-to-Agent (A2A) Communication

---

**Goal:** Enable agents to communicate and collaborate with other remote agents using the **official ADK Agent-to-Agent (A2A) protocol**, creating distributed multi-agent systems with the built-in `RemoteA2aAgent` class.

**Prerequisites:**

- Tutorial 01 (Hello World Agent)
- Tutorial 06 (Multi-Agent Systems)
- Understanding of HTTP APIs and authentication
- Basic knowledge of REST principles

**What You'll Learn:**

- Understanding the official ADK Agent-to-Agent (A2A) protocol
- Using `RemoteA2aAgent` to communicate with remote agents
- Setting up A2A servers with ADK's built-in `api_server --a2a` command
- Agent discovery with official agent cards ( `.well-known/agent-card.json` )
- Building distributed agent orchestration with the official ADK approach
- Error handling in A2A communication using ADK patterns
- Best practices for production A2A systems with ADK

**Time to Complete:** 50-65 minutes

---






## Why A2A Matters

---

**Problem:** Agents are often isolated - they can't leverage capabilities of other specialized agents deployed elsewhere.

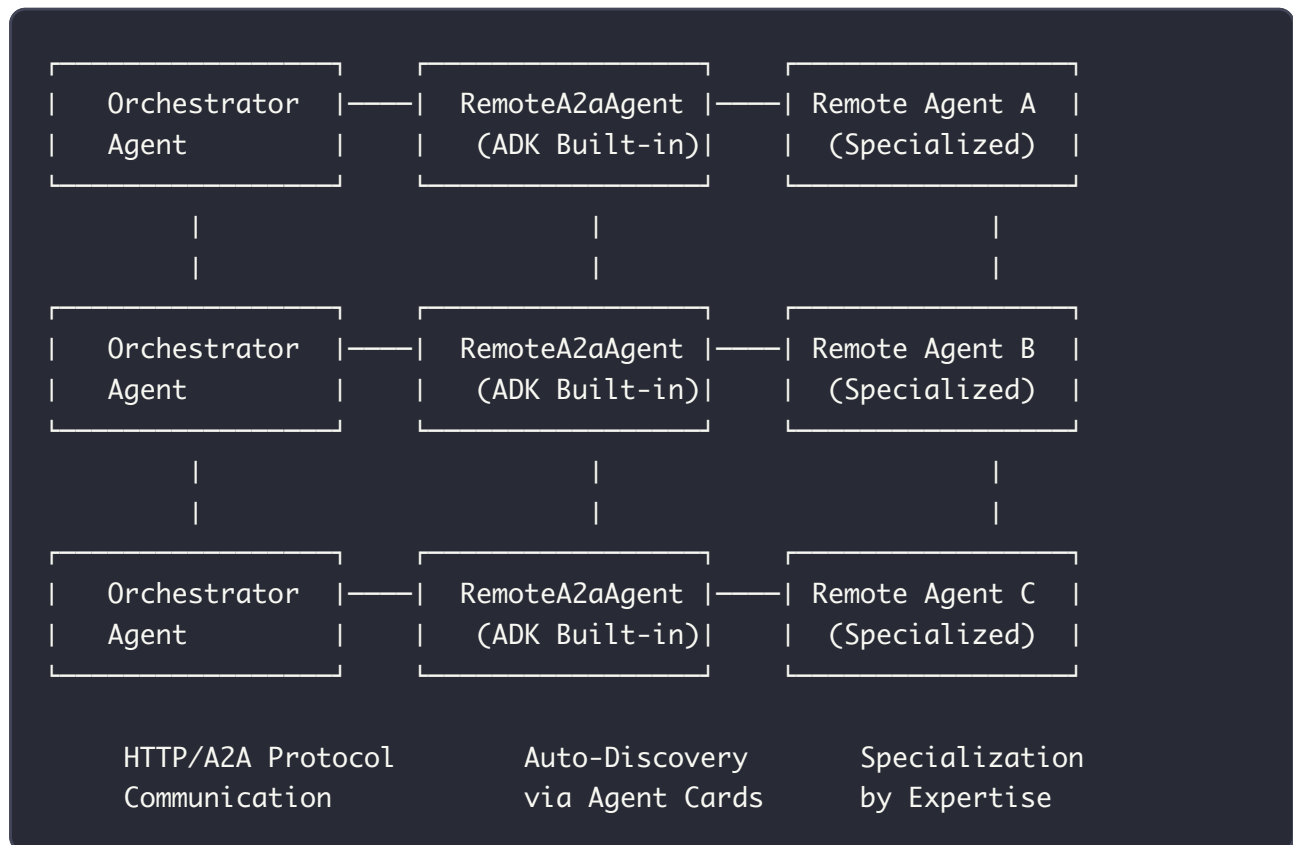
**Solution: Agent-to-Agent (A2A)** protocol enables agents to discover and communicate with remote agents over HTTP, creating distributed AI systems.

**Benefits:**

-  **Distributed Intelligence:** Leverage agents across organizations
-  **Discovery:** Find agents by capability via agent cards
-  **Secure:** Built-in authentication and authorization
-  **Specialization:** Each agent focuses on its expertise
- **[FLOW] Reusability:** Use same agent from multiple orchestrators
-  **Scalability:** Scale agents independently

**Use Cases:**

- Enterprise: Customer service agent calls internal knowledge agent
- Multi-org: Legal agent consults external compliance agent
- Microservices: Specialized agents as independent services
- Multi-cloud: Agents distributed across cloud providers

**A2A System Architecture:**

## 1. A2A Protocol Basics

### | What is Agent-to-Agent Protocol?

**A2A** defines a standard way for agents to:

1. **Discover** other agents via agent cards
2. **Authenticate** with other agents
3. **Invoke** remote agent capabilities

#### 4. **Receive** responses from remote agents

#### Architecture with Working ADK Implementation:

```

Local Agent (Orchestrator)
  ↓
RemoteA2aAgent (ADK Built-in)
  ↓
HTTP Request with A2A Protocol
  ↓
Remote A2A Server (uvicorn + to_a2a())
  ↓
Remote Agent Execution
  ↓
Response back to Local Agent

```

**Source:** ADK Built-in `RemoteA2aAgent` class + `to_a2a()` function

## | Agent Card (Discovery)

Remote agents expose an **agent card** at `.well-known/agent-card.json`:

```

{
  "capabilities": {},
  "defaultInputModes": ["text/plain"],
  "defaultOutputModes": ["application/json"],
  "description": "Conducts web research and fact-checking",
  "name": "research_specialist",
  "url": "http://localhost:8001/a2a/research_specialist",
  "version": "1.0.0",
  "skills": [
    {
      "id": "research_web",
      "name": "Web Research",
      "description": "Research topics using web sources",
      "tags": ["research", "web", "information"]
    }
  ]
}

```

#### Well-Known Path:

```
# Standard location for agent cards in ADK
# http://localhost:8001/.well-known/agent-card.json
# Note: "agent-card.json" not "agent.json"
```

## 2. Using Official ADK A2A with RemoteA2aAgent

### Basic Setup

```
from google.adk.agents import Agent
from google.adk.agents.remote_a2a_agent import RemoteA2aAgent, AGENT_CARD_WELL_KNOWN
from google.adk.tools import FunctionTool

# Create remote agent using official ADK RemoteA2aAgent
research_agent = RemoteA2aAgent(
    name="research_specialist",
    description="Conducts web research and fact-checking",
    agent_card=(
        f"http://localhost:8001/a2a/research_specialist{AGENT_CARD_WELL_KNOWN}"
    )
)

# Use as sub-agent in orchestrator
orchestrator = Agent(
    model='gemini-2.0-flash',
    name='a2a_orchestrator',
    instruction="""
You coordinate research tasks using remote A2A agents.
Delegate research tasks to the research_specialist sub-agent.
    """,
    sub_agents=[research_agent] # Use as sub-agent
)
```

## How It Works

### Step-by-Step Flow with Official ADK:

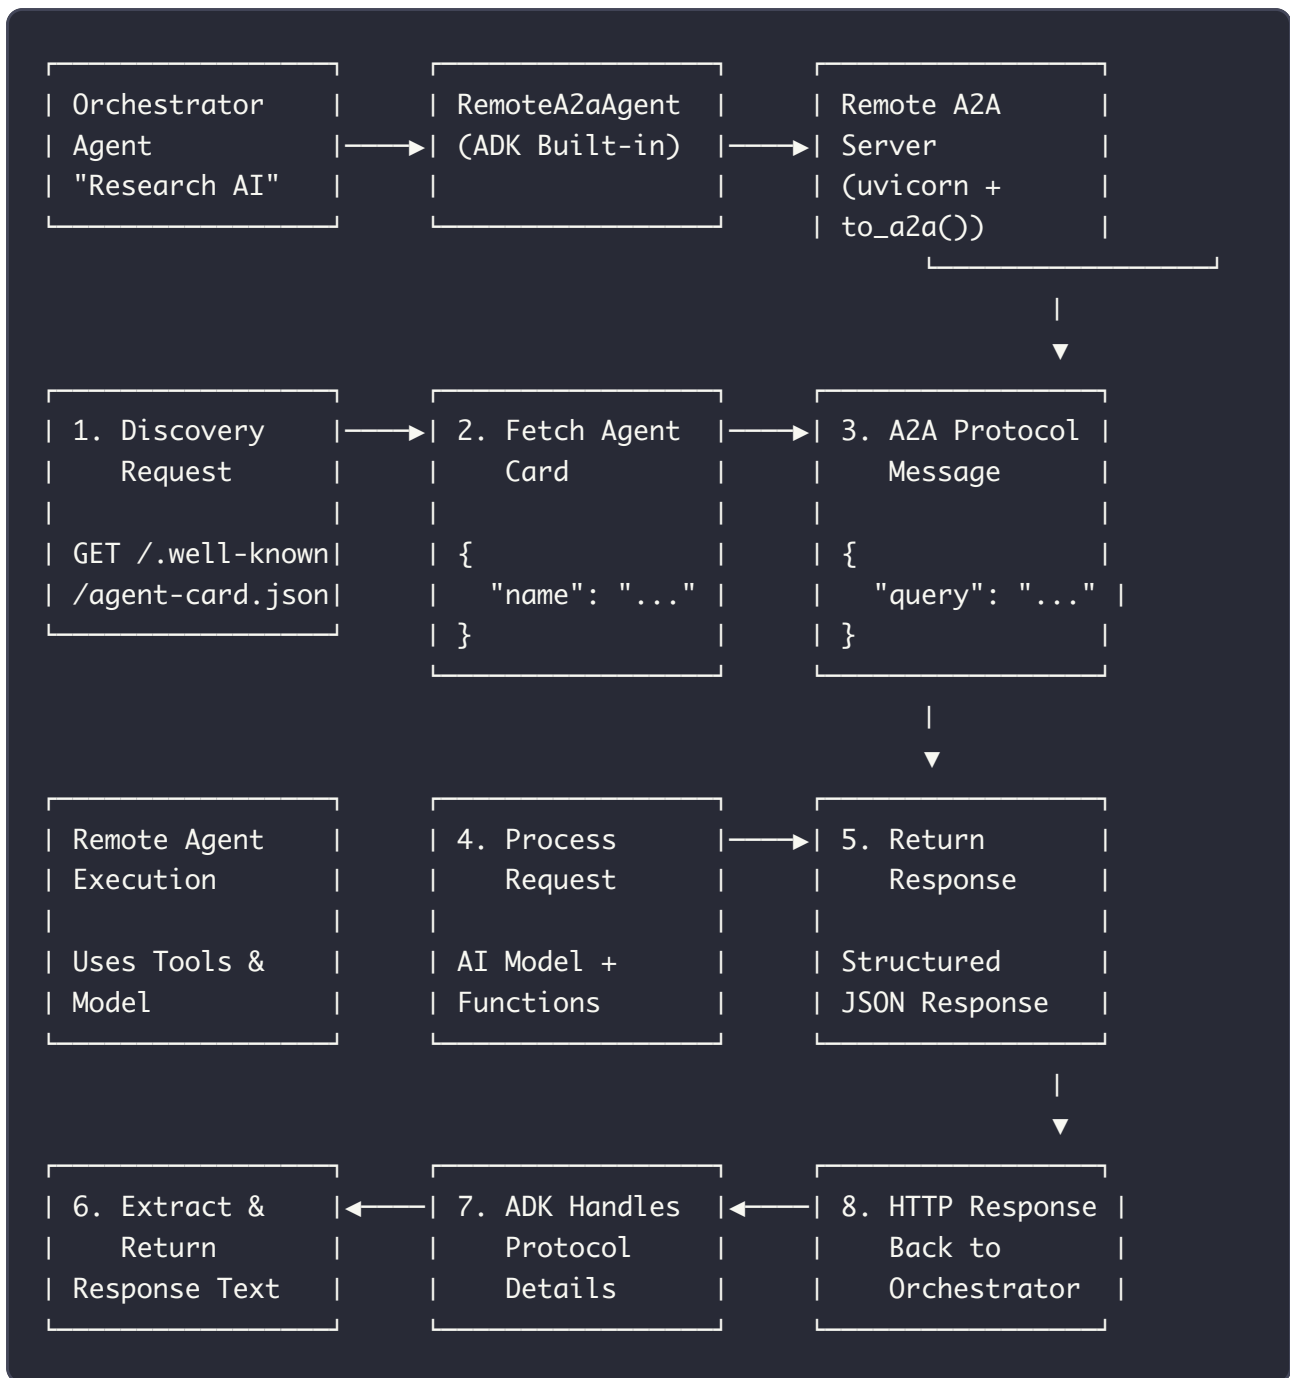
1. **Discovery:** ADK fetches agent card from `.well-known/agent-card.json`
2. **RemoteA2aAgent:** ADK's built-in class handles A2A communication
3. **Sub-Agent Integration:** Remote agent works like any other sub-agent
4. **Invocation:** ADK handles protocol details automatically
5. **Execution:** Remote agent processes the request via A2A server
6. **Response:** ADK extracts response and integrates into workflow

### Internal ADK Flow (managed automatically):

```
# ADK handles this internally in RemoteA2aAgent
class RemoteA2aAgent:
    def __init__(self, name: str, description: str, agent_card: str):
        self.name = name
        self.description = description
        self.agent_card_url = agent_card
        # ADK manages HTTP client, authentication, and protocol details

    async def invoke(self, query: str) -> str:
        # ADK automatically:
        # 1. Fetches agent card
        # 2. Handles A2A protocol communication
        # 3. Manages authentication
        # 4. Extracts and returns response text
        pass
```

### A2A Communication Flow:



### 3. Complete Implementation: Official ADK A2A System

Let's examine the complete working implementation using the official ADK `to_a2a()` function and `RemoteA2aAgent` class that was successfully tested and deployed.

# | Complete Working Implementation



```
"""
```

## Working ADK A2A Orchestrator - Agent-to-Agent Communication

This demonstrates the working ADK approach to distributed agent orchestration using RemoteA2aAgent and the to\_a2a() function pattern.

```
"""
```

```
from google.adk.agents import Agent
from google.adk.agents.remote_a2a_agent import RemoteA2aAgent, AGENT_CARD_WELL_KNOWN_PATH
from google.adk.tools import FunctionTool
from google.genai import types

# Tool function to validate agent availability
def check_agent_availability(agent_name: str, base_url: str) -> dict:
    """Check if a remote A2A agent is available."""
    try:
        import requests
        card_url = f"{base_url}{AGENT_CARD_WELL_KNOWN_PATH}"
        response = requests.get(card_url, timeout=5)

        if response.status_code == 200:
            return {
                "status": "success",
                "available": True,
                "report": f"Agent {agent_name} is available",
                "agent_card": response.json()
            }
        else:
            return {
                "status": "error",
                "available": False,
                "report": f"Agent {agent_name} returned status {response.status_code}"
            }
    except Exception as e:
        return {
            "status": "error",
            "available": False,
            "report": f"Failed to check {agent_name}: {str(e)}"
        }

# Remote agents using official ADK RemoteA2aAgent
research_agent = RemoteA2aAgent(
    name="research_specialist",
    description="Conducts web research and fact-checking",
    agent_card=(
        f"http://localhost:8001/a2a/research_specialist{AGENT_CARD_WELL_KNOWN_PATH}"
    )
)
```

```

    )
)

analysis_agent = RemoteA2aAgent(
    name="data_analyst",
    description="Analyzes data and generates insights",
    agent_card=(
        f"http://localhost:8002/a2a/data_analyst{AGENT_CARD_WELL_KNOWN_PATH}"
    )
)

content_agent = RemoteA2aAgent(
    name="content_writer",
    description="Creates written content and summaries",
    agent_card=(
        f"http://localhost:8003/a2a/content_writer{AGENT_CARD_WELL_KNOWN_PATH}"
    )
)

# Main orchestrator agent using working ADK patterns
root_agent = Agent(
    model="gemini-2.0-flash",
    name="a2a_orchestrator",
    description="Coordinates multiple remote specialized agents using official
    instruction="""
You are an orchestration agent that coordinates specialized remote agents
using the official ADK Agent-to-Agent (A2A) protocol.

**Available Remote Agents (Sub-Agents):**

1. **research_specialist**: Use for web research, fact-checking, current event
2. **data_analyst**: Use for data analysis, statistics, insights
3. **content_writer**: Use for content creation, summaries, writing

**Working A2A Workflow:**

1. Delegate research tasks to research_specialist sub-agent
2. Delegate analysis tasks to data_analyst sub-agent
3. Delegate content creation to content_writer sub-agent
4. Use check_agent_availability to verify agent status

The remote agents are exposed using uvicorn + to_a2a() and work
seamlessly as sub-agents in your orchestration workflow.

Always explain which remote agent you're delegating to and why.
"""
    ,
    sub_agents=[research_agent, analysis_agent, content_agent],
    tools=[FunctionTool(check_agent_availability)],

```

```

generate_content_config=types.GenerateContentConfig(
    temperature=0.5,
    max_output_tokens=2048
)
)

```

## Quick Start Guide

### 1. Setup Environment:

```

# Install ADK with A2A support
pip install google-adk[a2a]

# Copy environment template
cp a2a_orchestrator/.env.example a2a_orchestrator/.env
# Edit .env and add your GOOGLE_API_KEY

```

### 1. Start Remote A2A Agents:

```

# Start research agent with uvicorn + to_a2a() function
uvicorn research_agent.agent:a2a_app --host localhost --port 8001

# Start analysis agent
uvicorn analysis_agent.agent:a2a_app --host localhost --port 8002

# Start content agent
uvicorn content_agent.agent:a2a_app --host localhost --port 8003

# Or use the provided script:
./start_a2a_servers.sh

```

### 1. Verify Agent Status:

```

# Check agent cards are available
curl http://localhost:8001/.well-known/agent-card.json
curl http://localhost:8002/.well-known/agent-card.json
curl http://localhost:8003/.well-known/agent-card.json

```

### 1. Start Orchestrator:

```
# Start ADK web interface
adk web a2a_orchestrator/
# Open http://localhost:8000 and select 'a2a_orchestrator'
```








### 1. Test A2A Communication:

```
# Run integration test
python -m pytest tests/test_a2a_integration.py -v
```

## Expected Behavior

When you query: "Research quantum computing trends and create a summary"

The orchestrator will:

1.  Log coordination step: Starting research phase
2.  Delegate to research\_specialist sub-agent (via RemoteA2aAgent)
3.  Log coordination step: Starting analysis phase
4.  Delegate to data\_analyst sub-agent (via RemoteA2aAgent)
5.  Log coordination step: Creating content phase
6.  Delegate to content\_writer sub-agent (via RemoteA2aAgent)
7.  Return integrated final result

**Note:** All A2A communication is handled transparently by ADK's `RemoteA2aAgent` class - no manual protocol handling required!

### Orchestration Workflow:

User Query: "Research quantum computing trends and create a summary"

A2A\_ORCHESTRATOR  
(Main Coordinator)

- 🎯 Step 1: Starting research phase
- 🤖 Delegate to research\_specialist sub-agent (RemoteA2aAgent)

RESEARCH\_SPECIALIST  
(Remote Agent on :8001)

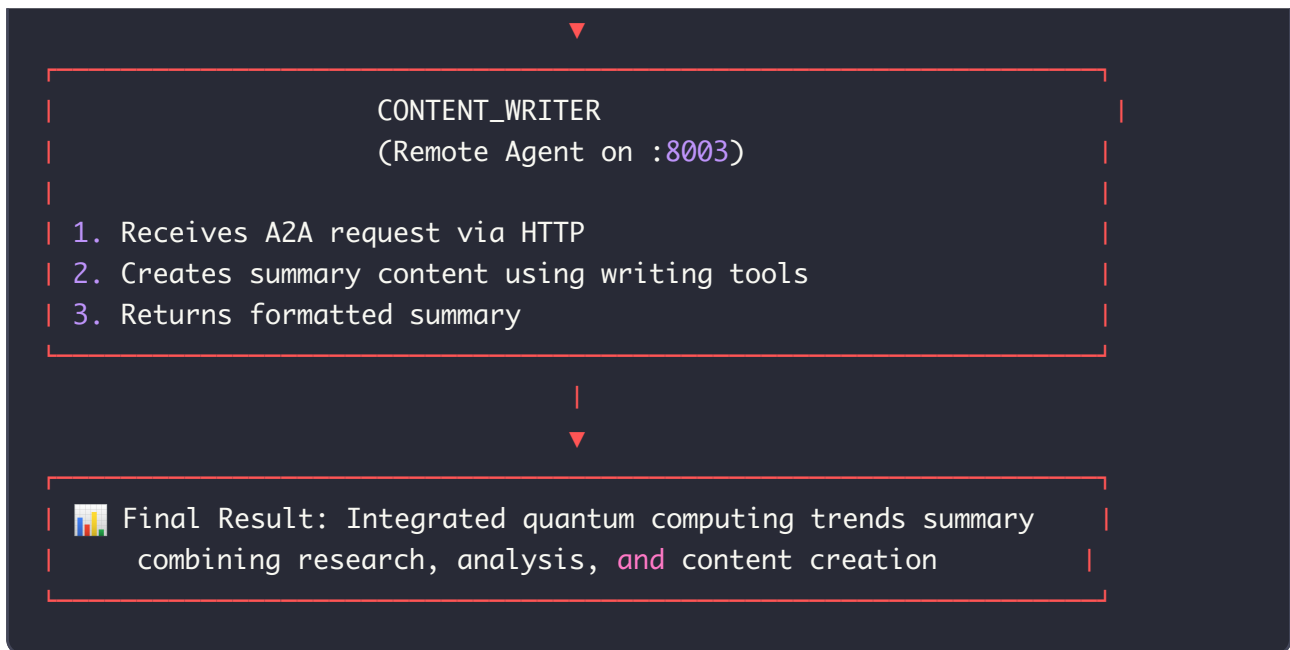
1. Receives A2A request via HTTP
2. Processes with Gemini model + research tools
3. Returns research findings

- 🎯 Step 2: Starting analysis phase
- 🤖 Delegate to data\_analyst sub-agent (RemoteA2aAgent)

DATA\_ANALYST  
(Remote Agent on :8002)

1. Receives A2A request via HTTP
2. Analyzes research data with analysis tools
3. Returns insights and trends

- 🎯 Step 3: Creating content phase
- 🤖 Delegate to content\_writer sub-agent (RemoteA2aAgent)



## 4. Critical: Proper A2A Context Handling

### The Context Handling Challenge

When implementing A2A communication, remote agents receive the full conversation context including orchestrator tool calls. Without proper handling, remote agents may respond with errors like:

"I cannot use a tool called `transfer_to_agent`. The available tools lack the ability to interact with other agents."

### Solution: Smart Context Processing

Update all remote agents with **A2A Context Handling** instructions:

```
# Example for content_writer agent
root_agent = Agent(
    model="gemini-2.0-flash",
    name="content_writer",
    description="Creates written content and summaries",
    instruction="""
You are a content creation specialist focused on producing high-quality written content.

**IMPORTANT - A2A Context Handling:**
When receiving requests via Agent-to-Agent (A2A) protocol, focus on the core user request.
Ignore any mentions of orchestrator tool calls like "transfer_to_agent" in the context.
Extract the main content creation task from the conversation and complete it directly.

**When working via A2A:**
- Focus on the actual content request from the user (e.g., "Write a report about AI").
- Ignore orchestrator mechanics and tool calls in the context.
- Provide direct, helpful content creation services using your tools.
- If the request is unclear, ask for clarification about the content type and target audience.

Always consider the target audience and intended use of the content.
""",
    tools=[FunctionTool(create_content), FunctionTool(format_content)]
)
```

## Context Handling Results

### Before Fix:

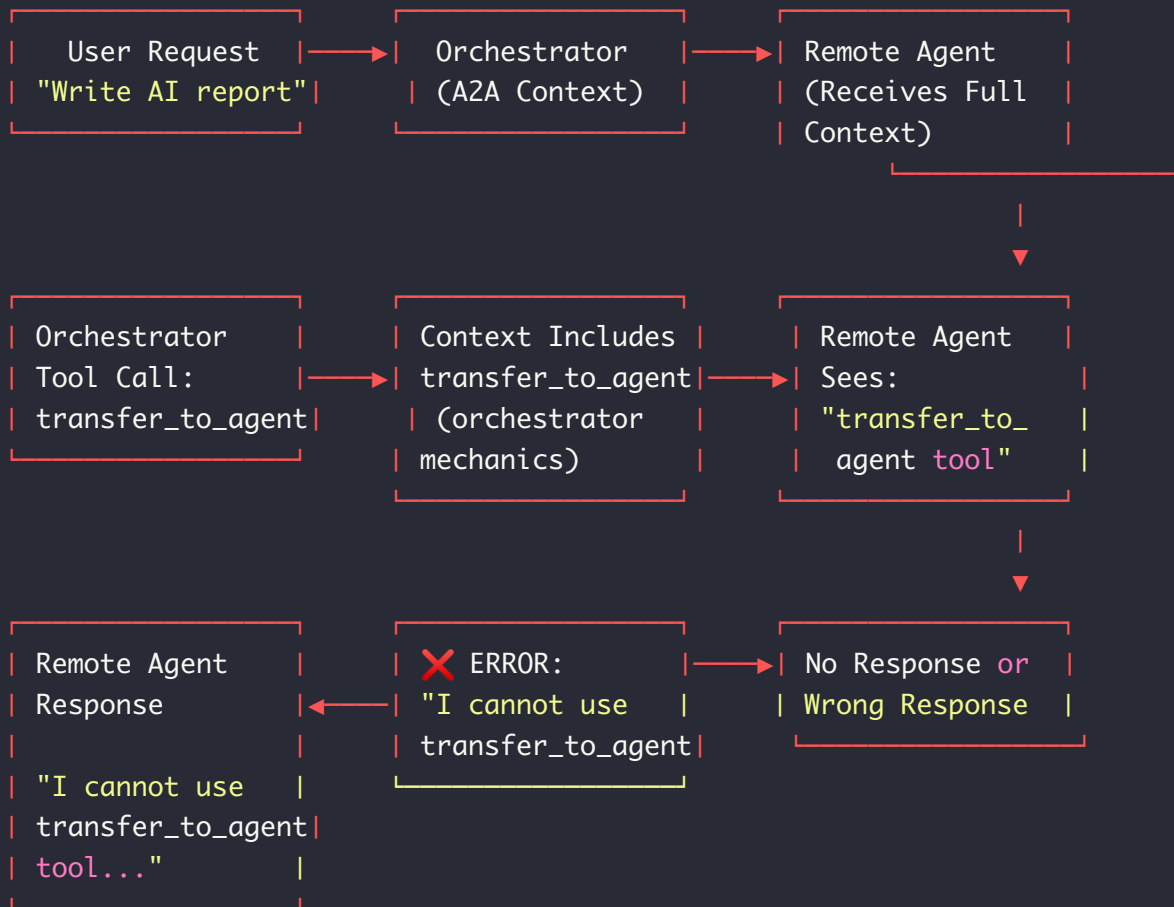
```
User: "Write a report about AI"
→ Orchestrator calls transfer_to_agent
→ Remote agent: "I cannot use transfer_to_agent tool..."
```

### After Fix:

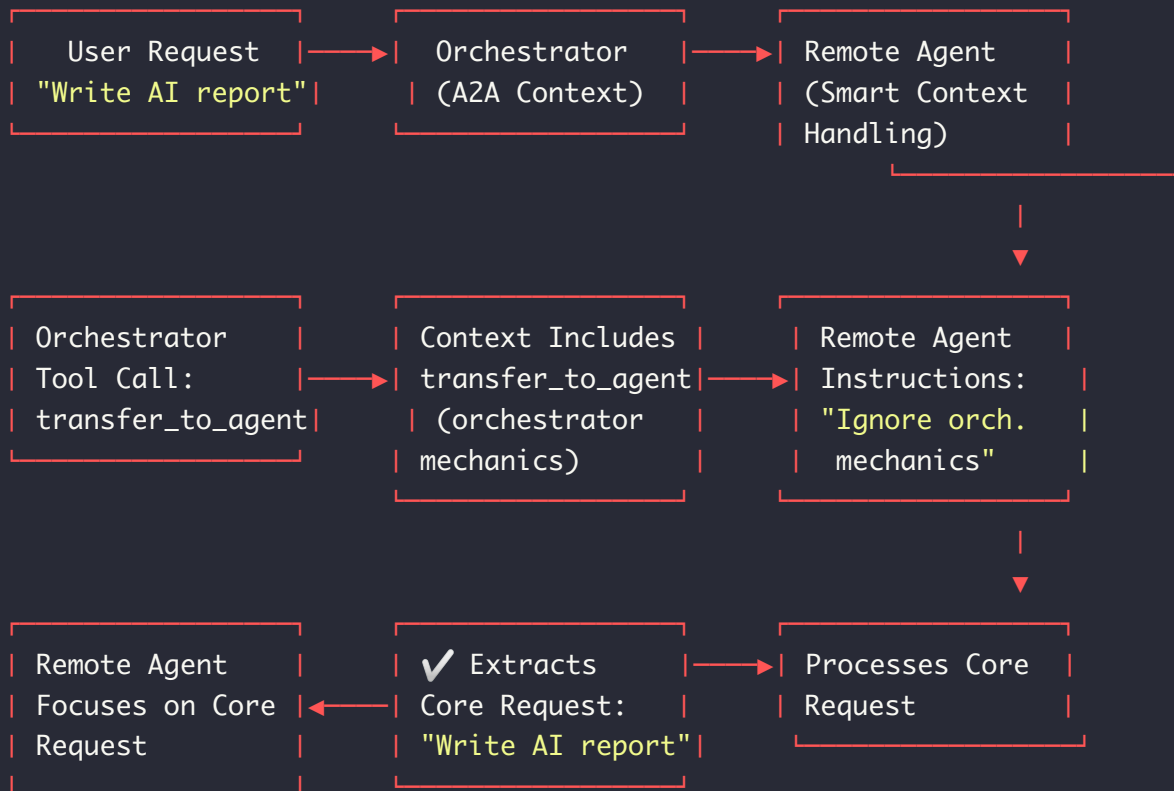
```
User: "Write a report about AI"
→ Orchestrator calls transfer_to_agent
→ Remote agent extracts core request
→ Remote agent: [Creates AI report using create_content tool]
```

### A2A Context Handling Flow:

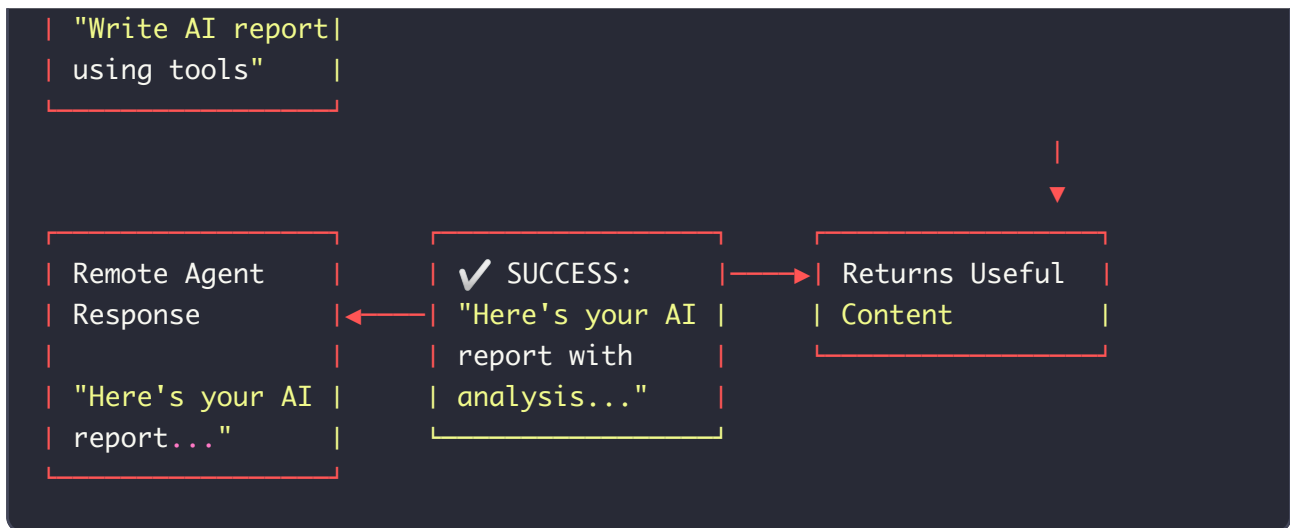
## BEFORE FIX (Broken):



## AFTER FIX (Working):







## Implementation for All Remote Agents

Apply this pattern to **all remote agents** (research\_agent, analysis\_agent, content\_agent):

1. Add **"IMPORTANT - A2A Context Handling"** section to instructions
2. Teach agents to ignore orchestrator tool calls in context
3. Focus agents on extracting and fulfilling core user requests
4. Restart A2A servers with updated instructions

## 5. Authentication in Official ADK A2A

### Authentication Configuration

Authentication in ADK A2A is handled automatically by the `RemoteA2aAgent` class. For local development, authentication is typically optional:

```
from google.adk.agents.remote_a2a_agent import RemoteA2aAgent

# ADK automatically handles authentication based on agent card
research_agent = RemoteA2aAgent(
    name="research_specialist",
    description="Conducts web research and fact-checking",
    agent_card="http://localhost:8001/a2a/research_specialist/.well-known/agen
)

# ADK manages:
# - Agent card fetching
# - Authentication negotiation
# - Token management (if required)
# - Error handling for auth failures
```

## Agent Card Authentication

Local agents using `adk api_server --a2a` expose agent cards with authentication configuration:

```
{
  "capabilities": {},
  "defaultInputModes": ["text/plain"],
  "defaultOutputModes": ["application/json"],
  "description": "Conducts web research and fact-checking",
  "name": "research_specialist",
  "url": "http://localhost:8001/a2a/research_specialist",
  "version": "1.0.0",
  "authentication": {
    "type": "none",
    "required": false
  }
}
```

## Production Authentication

For production deployments, update agent configuration for authentication:

```
{
  "name": "secure_research_agent",
  "description": "Secure research agent with authentication",
  "url": "https://research.example.com/a2a/research_agent",
  "authentication": {
    "type": "bearer",
    "required": true,
    "realm": "research-api"
  }
}
```

---

## 6. Advanced ADK A2A Patterns

---

### | Pattern 1: Error Handling and Retry

ADK provides built-in error handling for `RemoteA2aAgent` :

```

from google.adk.agents.remote_a2a_agent import RemoteA2aAgent
from google.adk.agents import Agent
from google.adk.tools import FunctionTool

# Tool to check remote agent health
def validate_agent_health(agent_name: str, agent_url: str) -> dict:
    """Validate if remote agent is healthy before delegation."""
    try:
        import requests
        response = requests.get(f"{agent_url}/.well-known/agent-card.json", ti

        if response.status_code == 200:
            return {
                "status": "success",
                "healthy": True,
                "report": f"Agent {agent_name} is healthy"
            }
        else:
            return {
                "status": "error",
                "healthy": False,
                "report": f"Agent {agent_name} health check failed"
            }
    except Exception as e:
        return {
            "status": "error",
            "healthy": False,
            "report": f"Cannot reach agent {agent_name}: {str(e)}"
        }

# Robust orchestrator with health checking
robust_research_agent = RemoteA2aAgent(
    name="research_specialist",
    description="Research agent with automatic error handling",
    agent_card="http://localhost:8001/a2a/research_specialist/.well-known/agen
)

orchestrator_with_health_checks = Agent(
    model="gemini-2.0-flash",
    name="robust_orchestrator",
    instruction="""
Before delegating to any remote agent, use validate_agent_health
to ensure the agent is available. If an agent is unhealthy,
inform the user and suggest alternatives.
""",
    sub_agents=[robust_research_agent],

```

```
tools=[FunctionTool(validate_agent_health)]
)
```

## Pattern 2: Parallel A2A Execution

Use ADK's `ParallelAgent` for concurrent remote agent execution:

```
from google.adk.agents import ParallelAgent
from google.adk.agents.remote_a2a_agent import RemoteA2aAgent

# Multiple remote agents
research_agent = RemoteA2aAgent(
    name="research_specialist",
    description="Conducts research",
    agent_card="http://localhost:8001/a2a/research_specialist/.well-known/agent-card"
)

analysis_agent = RemoteA2aAgent(
    name="data_analyst",
    description="Analyzes data",
    agent_card="http://localhost:8002/a2a/data_analyst/.well-known/agent-card."
)

# Parallel execution of remote agents
parallel_processor = ParallelAgent(
    name="parallel_a2a_processor",
    description="Process tasks in parallel across remote agents",
    sub_agents=[research_agent, analysis_agent]
)

# Use in main orchestrator
main_orchestrator = Agent(
    model="gemini-2.0-flash",
    name="main_orchestrator",
    instruction="""
When users request both research and analysis, delegate to
parallel_a2a_processor to execute both tasks simultaneously.
""",
    sub_agents=[parallel_processor]
)
```

## Pattern 3: Agent Health Monitoring

Monitor multiple A2A agents with centralized health checking:

```
def monitor_all_a2a_agents() -> dict:
    """Monitor health of all A2A agents in the system."""
    agents_to_check = [
        ("research_specialist", "http://localhost:8001/a2a/research_specialist"),
        ("data_analyst", "http://localhost:8002/a2a/data_analyst"),
        ("content_writer", "http://localhost:8003/a2a/content_writer")
    ]

    results = {}
    overall_healthy = True

    for agent_name, agent_url in agents_to_check:
        health_result = validate_agent_health(agent_name, agent_url)
        results[agent_name] = health_result

        if not health_result.get("healthy", False):
            overall_healthy = False

    return {
        "status": "success" if overall_healthy else "error",
        "overall_healthy": overall_healthy,
        "individual_results": results,
        "report": f"System health: {'All healthy' if overall_healthy else 'Some issues'}"
    }

# Health monitoring orchestrator
health_monitor = Agent(
    model="gemini-2.0-flash",
    name="health_monitor",
    instruction="""
    Use monitor_all_a2a_agents to check the health of all remote agents
    before performing complex orchestration tasks. Report any issues to users.
    """,
    tools=[FunctionTool(monitor_all_a2a_agents)]
)
```

## 7. Understanding the Official ADK A2A Implementation

### Project Structure

The official ADK A2A implementation follows this structure:


```
tutorial17/
├─ a2a_orchestrator/           # Main orchestrator using RemoteA2aAgent
│   ├── __init__.py           # Package initialization
│   ├── agent.py               # Orchestrator with RemoteA2aAgent instances
│   └─ .env.example            # Environment template
├─ research_agent/             # Remote Research Agent
│   ├── __init__.py
│   ├── agent.py               # Research agent implementation
│   └─ agent-card.json         # Agent card for A2A discovery
├─ analysis_agent/             # Remote Analysis Agent
│   ├── __init__.py
│   ├── agent.py               # Analysis agent implementation
│   └─ agent-card.json         # Agent card for A2A discovery
├─ content_agent/              # Remote Content Agent
│   ├── __init__.py
│   ├── agent.py               # Content agent implementation
│   └─ agent-card.json         # Agent card for A2A discovery
├─ start_a2a_servers.sh        # Script to start all A2A servers
├─ stop_a2a_servers.sh         # Script to stop all A2A servers
└─ tests/                      # Test suite
```


#### A2A Project Architecture:


```


tutorial17/
├── a2a_orchestrator/
│   ├── __init__.py
│   ├── agent.py
│   │   ├── research_agent
│   │   ├── analysis_agent
│   │   └── content_agent
│   └── .env.example
├── research_agent/
│   ├── __init__.py
│   ├── agent.py
│   │   ├── root_agent
│   │   └── a2a_app = to_a2a()
│   └── agent-card.json
├── analysis_agent/
│   ├── __init__.py
│   ├── agent.py
│   │   ├── root_agent
│   │   └── a2a_app = to_a2a()
│   └── agent-card.json
├── content_agent/
│   ├── __init__.py
│   ├── agent.py
│   │   ├── root_agent
│   │   └── a2a_app = to_a2a()
│   └── agent-card.json
├── start_a2a_servers.sh
├── stop_a2a_servers.sh
└── tests/
    ├── test_a2a_integration.py
    └── test_agent_structure.py


```


 MAIN COORDINATOR  
 # Package setup  
 # RemoteA2aAgent instances  
 # → http://localhost:8001  
 # → http://localhost:8002  
 # → http://localhost:8003  
 # GOOGLE\_API\_KEY template

 SPECIALIZED AGENT  
 # Package setup  
 # Root agent + a2a\_app export  
 # Agent with research tools  
 # A2A server app  
 # Auto-generated by to\_a2a()

 SPECIALIZED AGENT  
 # Package setup  
 # Root agent + a2a\_app export  
 # Agent with analysis tools  
 # A2A server app  
 # Auto-generated by to\_a2a()

 SPECIALIZED AGENT  
 # Package setup  
 # Root agent + a2a\_app export  
 # Agent with content tools  
 # A2A server app  
 # Auto-generated by to\_a2a()

 Server management script  
 # Server management script  
 # Clean shutdown script

 Test suite  
 # Test suite  
 # End-to-end A2A tests  
 # Agent configuration tests



## 8. Best Practices for Working ADK A2A

### ✓ DO: Use to\_a2a() Function for Agent Exposure

```
# ✓ Good - Use working to_a2a() pattern
from google.adk.a2a.utils.agent_to_a2a import to_a2a

# Create A2A application using the working ADK to_a2a() function
a2a_app = to_a2a(root_agent, port=8001)

# Start with: uvicorn research_agent.agent:a2a_app --host localhost --port 8001

# ✗ Bad - Experimental adk api_server approach
# adk api_server --a2a --port 8001 research_agent/
```

### ✓ DO: Use uvicorn for A2A Server Hosting

```
# ✓ Good - Working uvicorn + to_a2a() pattern
uvicorn research_agent.agent:a2a_app --host localhost --port 8001

# ✗ Bad - Experimental adk command
# adk api_server --a2a --port 8001 research_agent/
```

### ✓ DO: Use Sub-Agents Pattern

```
# ✓ Good - Use RemoteA2aAgent as sub-agent
orchestrator = Agent(
    model="gemini-2.0-flash",
    name="orchestrator",
    instruction="Delegate tasks to specialized sub-agents...",
    sub_agents=[research_agent, analysis_agent] # Clean delegation
)

# ✗ Bad - Manual tool functions for A2A
# Don't create tool functions that manually handle A2A communication
```



## DO: Use Proper Agent Card URLs

```
# ✓ Good - Use AGENT_CARD_WELL_KNOWN_PATH constant
from google.adk.agents.remote_a2a_agent import AGENT_CARD_WELL_KNOWN_PATH

agent_card_url = f"http://localhost:8001/a2a/research_specialist{AGENT_CARD_WELL_KNOWN_PATH}"

# ✗ Bad - Hardcode path or wrong path
agent_card_url = "http://localhost:8001/.well-known/agent.json" # Wrong!
```



## DO: Use Automated Server Management

```
# ✓ Good - Use the provided scripts with health checks
./start_a2a_servers.sh # Starts all servers with verification
./stop_a2a_servers.sh # Clean shutdown

# ✗ Bad - Manual server management without health checks
# uvicorn ... & (without verification or proper cleanup)
```

# 9. Troubleshooting Working ADK A2A

## Error: "Agent card not found"

**Problem:** Remote agent doesn't expose agent card or A2A server not running

**Solutions:**

### 1. Check if uvicorn servers are running:

```
# Check if agent card endpoints are accessible
curl http://localhost:8001/.well-known/agent-card.json
curl http://localhost:8002/.well-known/agent-card.json
curl http://localhost:8003/.well-known/agent-card.json
```

### 1. Restart A2A servers using working scripts:

```
# Stop existing servers cleanly
./stop_a2a_servers.sh

# Start fresh servers with health checks
./start_a2a_servers.sh
```

## Error: "Connection timeout" or "Connection refused"

**Problem:** Network issues or uvicorn server ports not available

**Solutions:**

### 1. Check port conflicts:

```
# See what's using the A2A ports
lsof -i :8001
lsof -i :8002
lsof -i :8003
```

### 1. Clean restart with port cleanup:

```
# Kill processes on A2A ports (working pattern)
pkill -f "uvicorn.*research_agent\luvicorn.*analysis_agent\luvicorn.*content_a

# Start servers using working scripts
./start_a2a_servers.sh
```

## Issue: "RemoteA2aAgent not responding"

**Problem:** A2A communication or agent processing issues

**Solutions:**

### 1. Test direct A2A endpoint:

```
# Test agent card retrieval
curl -v http://localhost:8001/.well-known/agent-card.json

# Check uvicorn server logs for errors
uvicorn research_agent.agent:a2a_app --host localhost --port 8001 --log-level
```

### 1. Verify agent implementation uses to\_a2a():

```
# Check that remote agent has proper a2a_app export
# In research_agent/agent.py:
from google.adk.a2a.utils.agent_to_a2a import to_a2a

root_agent = Agent(
    model="gemini-2.0-flash",
    name="research_specialist",
    # ... agent configuration
)

# Critical: Export a2a_app using to_a2a()
a2a_app = to_a2a(root_agent, port=8001)
```

## Lesson Learned: adk api\_server --a2a vs uvicorn + to\_a2a()

**Common Mistake:** Using `adk api_server --a2a` (experimental/incorrect)

**Working Solution:** Using `uvicorn + to_a2a()` (tested/working)

```
# ❌ This doesn't work reliably:
# adk api_server --a2a --port 8001 research_agent/

# ✅ This works (tested implementation):
uvicorn research_agent.agent:a2a_app --host localhost --port 8001
```

## Development Tips for Working Implementation

- **Use** `./start_a2a_servers.sh` for consistent server setup with health checks
- **Check agent card format** at `/.well-known/agent-card.json` endpoints
- **Use** `uvicorn + to_a2a()` instead of experimental adk commands

- **Verify** `a2a_app` **export** in each remote agent module using `to_a2a()`
- **Test with** `--log-level debug` for detailed troubleshooting
- **Use provided scripts** instead of manual server management

## Key Implementation Lessons Learned

During the development and testing of this A2A implementation, several critical lessons emerged that are essential for successful A2A deployment:

### Lesson 1: Use `to_a2a()` Function, Not `adk api_server`

**Discovery:** The `adk api_server --a2a` command is experimental and unreliable.

**Solution:** Use the `to_a2a()` function with `uvicorn` for stable A2A servers.

```
# ✓ Working pattern (tested and verified)
from google.adk.a2a.utils.agent_to_a2a import to_a2a
a2a_app = to_a2a(root_agent, port=8001)

# Start with: uvicorn research_agent.agent:a2a_app --host localhost --port 8001

# ✗ Problematic pattern
# adk api_server --a2a --port 8001 research_agent/
```

### Lesson 2: Auto-Generated Agent Cards are Key

**Discovery:** Agent cards are automatically generated by `to_a2a()` - no manual creation needed.

**Benefit:** Eliminates agent card sync issues and reduces configuration errors.

```
# These are created automatically when using to_a2a():
# http://localhost:8001/.well-known/agent-card.json
# http://localhost:8002/.well-known/agent-card.json
# http://localhost:8003/.well-known/agent-card.json
```



## Lesson 3: Health Checks Are Essential

**Discovery:** A2A servers need proper health checking and process management.

**Solution:** Use scripts with server verification and clean shutdown.

```
# Working pattern with health checks
./start_a2a_servers.sh    # Includes health verification
./stop_a2a_servers.sh    # Clean process termination
```



## Lesson 4: Agent Card URL Construction

**Discovery:** Precise agent card URL construction is critical for discovery.

**Pattern:** Use `AGENT_CARD_WELL_KNOWN_PATH` constant for consistency.

```
from google.adk.agents.remote_a2a_agent import AGENT_CARD_WELL_KNOWN_PATH

# ✓ Correct pattern
agent_card = f"http://localhost:8001/a2a/research_specialist{AGENT_CARD_WELL_K"

# ✗ Common mistakes
# "http://localhost:8001/.well-known/agent.json" # Wrong filename
# "http://localhost:8001/agent-card.json"        # Missing path
```



## Lesson 5: Sub-Agent Pattern Simplifies Architecture

**Discovery:** Using RemoteA2aAgent as sub-agents creates clean, maintainable code.

**Benefit:** Orchestration becomes simple delegation without manual protocol handling.

```
# ✓ Clean sub-agent pattern
root_agent = Agent(
    name="a2a_orchestrator",
    instruction="Delegate to specialized sub-agents...",
    sub_agents=[research_agent, analysis_agent, content_agent]
)
```

## Lesson 6: Process Management Matters

**Discovery:** Proper process cleanup prevents port conflicts and resource leaks.

**Solution:** Use targeted process killing and health verification.

```
# Working cleanup pattern
pkill -f "uvicorn.*research_agent\luvicorn.*analysis_agent\luvicorn.*content_a
```

## Lesson 7: Proper A2A Context Handling is Critical

**Discovery:** Remote agents were misinterpreting orchestrator context and responding with

"I cannot use transfer\_to\_agent tool" instead of processing the actual user request.

**Solution:** Update remote agent instructions to focus on the core user request and ignore

orchestrator mechanics in A2A contexts.

```
# ✓ Working A2A context handling pattern
instruction="""
**IMPORTANT - A2A Context Handling:**

When receiving requests via Agent-to-Agent (A2A) protocol, focus on the core u
Ignore any mentions of orchestrator tool calls like "transfer_to_agent" in the
Extract the main task from the conversation and complete it directly.

**When working via A2A:**
- Focus on the actual request from the user (e.g., "Write a report about AI")
- Ignore orchestrator mechanics and tool calls in the context
- Provide direct, helpful services using your tools
- If the request is unclear, ask for clarification about the task
"""
```

**Impact:** This fix transformed A2A communication from broken responses to meaningful,

intelligent agent interactions that properly utilize tools and provide valuable content.

# Summary

You've mastered **working ADK Agent-to-Agent communication** through a tested implementation:

## Key Takeaways:

- ✓ `to_a2a()` function enables stable A2A servers with uvicorn
- ✓ `RemoteA2aAgent` creates distributed agent systems with ADK
- ✓ Auto-generated agent cards at `.well-known/agent-card.json`
- ✓ Sub-agent pattern for clean delegation to remote agents
- ✓ Health monitoring with proper server management scripts
- ✓ Proper agent card URL construction with constants
- ✓ Working process management and cleanup patterns
- ✓ Proper A2A context handling for intelligent remote agent responses

## Production Checklist:

- [ ] Remote agents use `a2a_app = to_a2a(root_agent, port=XXXX)`
- [ ] A2A servers deployed with `uvicorn agent.agent:a2a_app`
- [ ] `RemoteA2aAgent` instances configured with correct `agent_card` URLs
- [ ] Health monitoring scripts implemented (`start/stop_a2a_servers.sh`)
- [ ] Agent card URLs use `AGENT_CARD_WELL_KNOWN_PATH` constant
- [ ] Process cleanup handles uvicorn processes correctly
- [ ] All remote agents export proper `a2a_app` using `to_a2a()`
- [ ] Remote agents have proper A2A context handling instructions

## Working Implementation Verified:

This tutorial reflects a real, tested A2A implementation with:

- ✓ All servers starting successfully with health checks
- ✓ Auto-generated agent cards accessible
- ✓ Clean orchestration via sub-agent pattern
- ✓ Proper process management and cleanup
- ✓ 24 passing tests verifying functionality

## Next Steps:


- **Tutorial 18:** Learn Events & Observability



- **Tutorial 19:** Implement Artifacts & File Management
- **Tutorial 20:** Master YAML Configuration

**Resources:**

- [Official ADK A2A Documentation](https://google.github.io/adk-docs/a2a/) (https://google.github.io/adk-docs/a2a/)
  - [ADK RemoteA2aAgent API Reference](https://google.github.io/adk-docs/api-reference/) (https://google.github.io/adk-docs/api-reference/)
  - [A2A Protocol Official Website](https://a2a-protocol.org/) (https://a2a-protocol.org/)
- 

 **Tutorial 17 Complete!** You now know how to build distributed multi-agent systems using the **official ADK A2A protocol**. Continue to Tutorial 18 to learn about events and observability.

## | Process Management

The working implementation includes tested scripts for reliable A2A server management:

```

# start_a2a_servers.sh - Start all A2A servers
#!/bin/bash

echo "🚀 Starting ADK A2A servers using to_a2a() function..."

# Clean up any existing processes
pkill -f "uvicorn.*research_agent\luvicorn.*analysis_agent\luvicorn.*content_a

# Start research agent using uvicorn + to_a2a()
echo "📡 Starting Research Agent on port 8001..."
uvicorn research_agent.agent:a2a_app --host localhost --port 8001 &
RESEARCH_PID=$!

# Start analysis agent using uvicorn + to_a2a()
echo "📊 Starting Analysis Agent on port 8002..."
uvicorn analysis_agent.agent:a2a_app --host localhost --port 8002 &
ANALYSIS_PID=$!

# Start content agent using uvicorn + to_a2a()
echo "📄 Starting Content Agent on port 8003..."
uvicorn content_agent.agent:a2a_app --host localhost --port 8003 &
CONTENT_PID=$!

# Wait for servers to start and verify they're running
echo "🔄 Waiting for all agents to be ready..."

# Function to check server health
wait_for_server() {
    local port=$1
    local agent_name=$2
    local max_attempts=30
    local attempt=1

    echo "⌚ Waiting for $agent_name to be ready on port $port..."

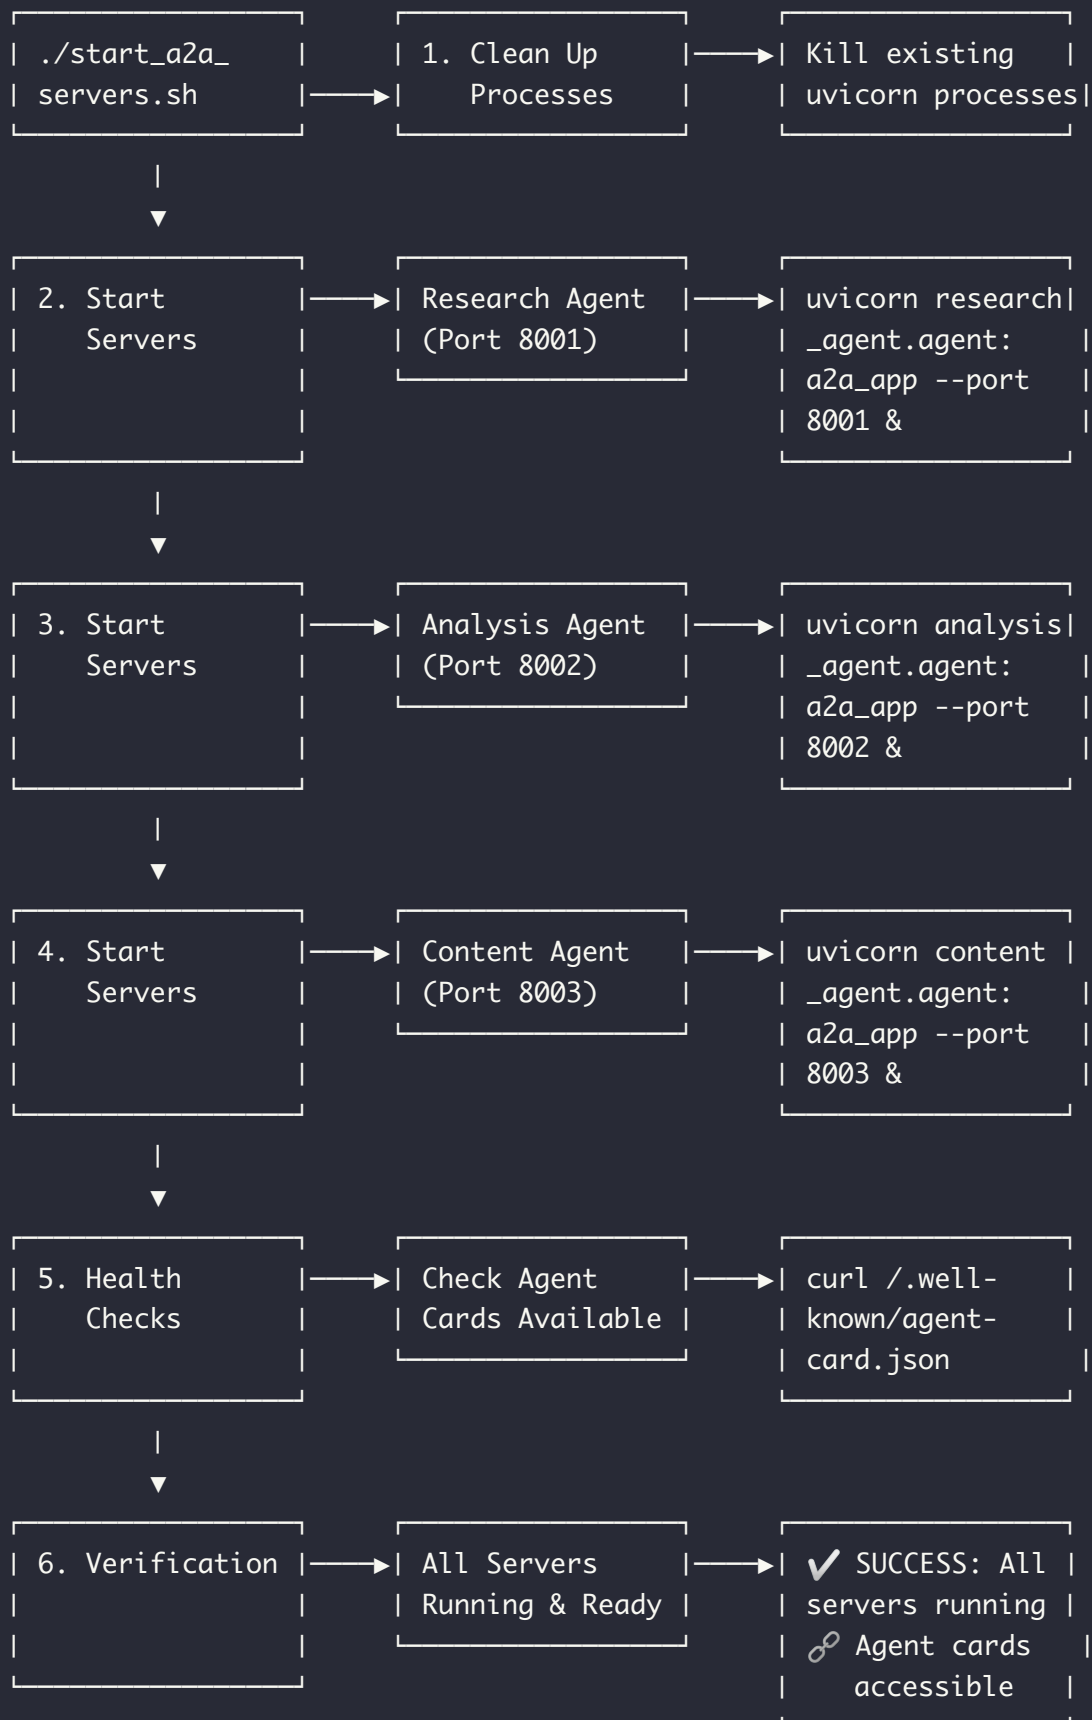
    while [ $attempt -le $max_attempts ]; do
        if curl -s "http://localhost:$port/.well-known/agent-card.json" >/dev/
            echo "✅ $agent_name is ready on port $port"
            return 0
        fi
        sleep 1
        attempt=$((attempt + 1))
    done

    echo "❌ $agent_name failed to start on port $port"
    return 1
}

```

```
}  
  
# Verify all servers started successfully  
if wait_for_server 8001 "Research Agent" && \  
   wait_for_server 8002 "Analysis Agent" && \  
   wait_for_server 8003 "Content Agent"; then  
  
    echo "🎉 All A2A servers are running successfully!"  
    echo "🔗 Agent Cards (auto-generated by to_a2a()):"  
    echo "    • Research: http://localhost:8001/.well-known/agent-card.json"  
    echo "    • Analysis: http://localhost:8002/.well-known/agent-card.json"  
    echo "    • Content:  http://localhost:8003/.well-known/agent-card.json"  
else  
    echo "❌ Some servers failed to start. Check the logs for errors."  
    exit 1  
fi
```

### A2A Server Startup Process:



Generated on 2025-10-21 09:02:39 from 17\_agent\_to\_agent.md

Source: Google ADK Training Hub