

# Tutorial 34: Google Cloud Pub/Sub + Event-Driven Agents

---

**Difficulty:** advanced

**Reading Time:** 1 hour

**Tags:** cloud, pubsub, event-driven, python, agents

**Description:** Build event-driven document processing pipelines with Google Cloud Pub/Sub and ADK agents for asynchronous processing.

This tutorial implements a real event-driven document processing system using Google Cloud Pub/Sub and ADK agents. It demonstrates a coordinator + specialist agents pattern with structured JSON output using Pydantic models.

Verified as of October 2025 with latest ADK and Gemini 2.5 Flash.

**Estimated Reading Time:** 50-60 minutes

**Difficulty Level:** Advanced

**Prerequisites:** Tutorial 01-03 (ADK Basics), Google Cloud project

---



## Quick Start - Working Implementation

---

The easiest way to get started is with our **complete working implementation**:

```
cd tutorial_implementation/tutorial34
make setup      # Install dependencies
make test       # Run all tests
```



[View Full Implementation](#) (../..../tutorial\_implementation/tutorial34)

### What's included:

- ✓ `root_agent`: Coordinator agent that routes documents to specialists
  - ✓ 4 Specialist agents: Financial, Technical, Sales, Marketing analyzers
  - ✓ Pydantic output schemas: Structured JSON results
  - ✓ 66 comprehensive tests (all passing)
  - ✓ Real-world example code ready to run
- 

## Table of Contents

---

1. [Overview](#) ([#overview](#))
  2. [Prerequisites & Setup](#) ([#prerequisites--setup](#))
  3. [Understanding the Architecture](#) ([#understanding-the-architecture](#))
  4. [Core Components](#) ([#core-components](#))
  5. [Running Locally](#) ([#running-locally](#))
  6. [Google Cloud Deployment](#) ([#google-cloud-deployment](#))
  7. [Troubleshooting](#) ([#troubleshooting](#))
  8. [Next Steps](#) ([#next-steps](#))
- 

## Overview

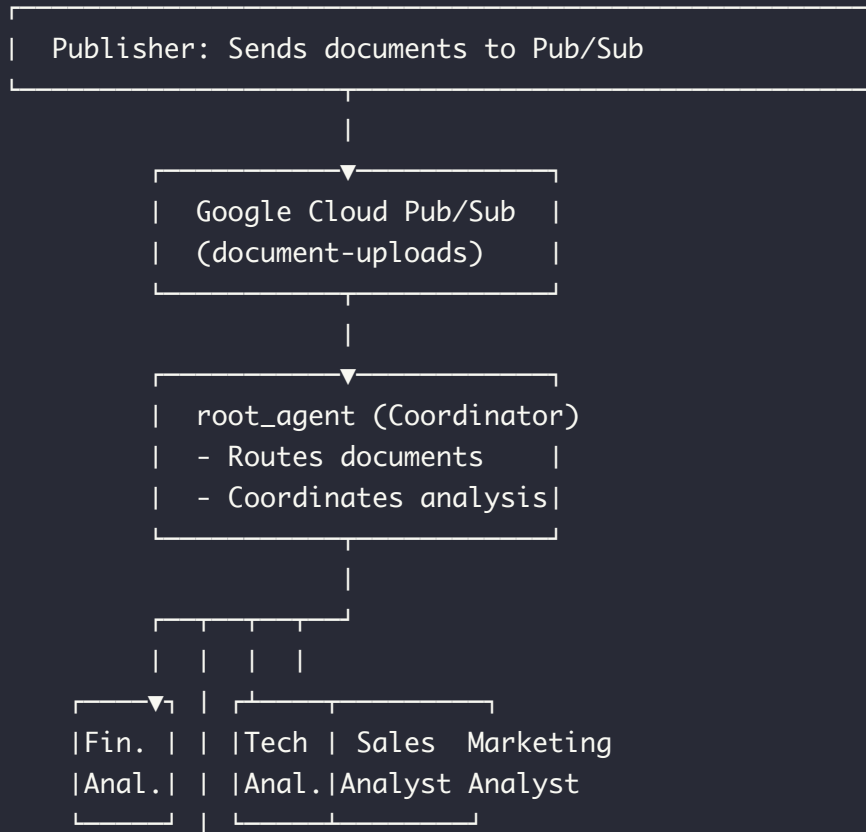
---

### | What You'll Build

In this tutorial, you'll build an **event-driven document processing system** using:

- **Google Cloud Pub/Sub** (Event messaging)
- **Google ADK** (Multi-agent coordination)
- **Gemini 2.5 Flash** (Document analysis)
- **Pydantic Models** (Structured JSON output)

### Architecture:



## Why Pub/Sub + ADK?

Feature	Benefit
<b>Asynchronous</b>	Non-blocking processing
<b>Decoupled</b>	Publishers and subscribers independent
<b>Scalable</b>	Auto-scales message volume
<b>Structured</b>	Pydantic models for JSON
<b>Reliable</b>	At-least-once delivery, retries

**When to use Pub/Sub + ADK:**

- ✓ Asynchronous document processing
  - ✓ Multi-step workflows
  - ✓ Event-driven architectures
  - ✓ Systems with strict output schemas
  - ✓ Google Cloud deployments
  - ✗ Real-time chat interfaces → Use Next.js/WebSocket
  - ✗ Simple synchronous calls → Use direct API
- 

## Prerequisites & Setup

---

### | Local Testing (No GCP Required)

To get started without Google Cloud:

```
# Install dependencies
cd tutorial_implementation/tutorial34
make setup

# Run tests - verifies agent configuration
make test

# This works completely locally using in-memory processing
```

### | Google Cloud Setup (Optional - For Real Pub/Sub)

To deploy with real Google Cloud Pub/Sub:

## 1. Install gcloud CLI

```
# macOS
brew install --cask google-cloud-sdk

# Then initialize
gcloud init
```

## 2. Authenticate

```
# Login to Google Cloud
gcloud auth login

# Set default project
gcloud config set project your-project-id

# Verify authentication
gcloud auth list
```

## 3. Create Pub/Sub Resources

```
# Enable Pub/Sub API
gcloud services enable pubsub.googleapis.com

# Create topic
gcloud pubsub topics create document-uploads

# Create subscription
gcloud pubsub subscriptions create document-processor \
  --topic=document-uploads \
  --ack-deadline=600
```

## 4. Set Environment Variables

```
# Set your GCP project

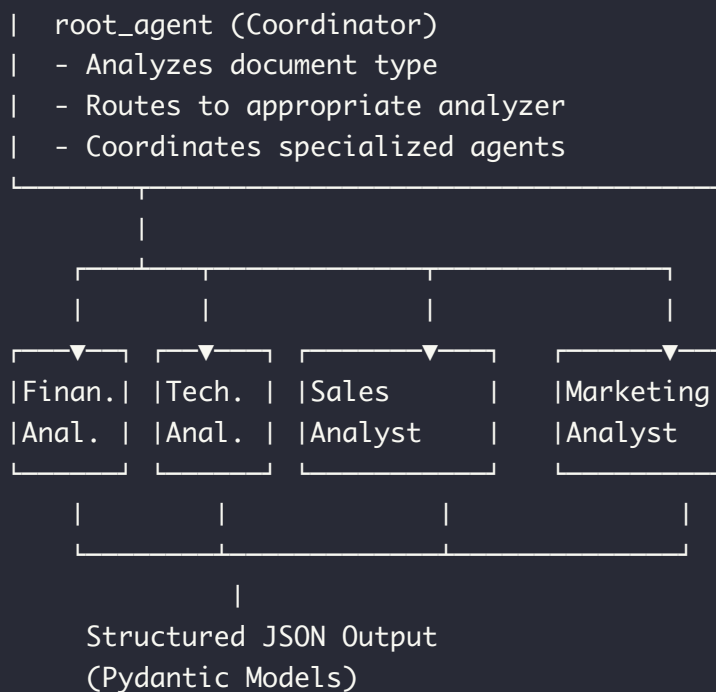
# Set Gemini API key

# Set application credentials
gcloud auth application-default login
```

## Understanding the Architecture

### The Coordinator + Specialist Pattern

This implementation uses a **coordinator agent** that intelligently routes documents to specialized analyzers:



## Key Components

1. **root\_agent** ( `pubsub_agent/agent.py` ):
2. Coordinator that routes documents to specialists
3. Analyzes document type and content
4. Calls appropriate sub-agent tool
5. Returns structured analysis
6. **Sub-Agents** (financial, technical, sales, marketing):
7. Specialized analyzers for document types
8. Enforce structured JSON via Pydantic `output_schema`
9. Extract type-specific metrics and insights
10. **Pydantic Output Schemas:**
11. `FinancialAnalysisOutput` : Revenue, profit, metrics
12. `TechnicalAnalysisOutput` : Technologies, components
13. `SalesAnalysisOutput` : Deals, pipeline value
14. `MarketingAnalysisOutput` : Campaigns, engagement metrics

## Pub/Sub Guarantees

Feature	Description
<b>At-least-once</b>	Messages delivered $\geq 1$ time
<b>Asynchronous</b>	Non-blocking processing
<b>Scalable</b>	Auto-scales message volume
<b>Durable</b>	Messages stored in topics
<b>Reliable</b>	Automatic retries on failure

# Core Components

## Agent Configuration

View the agent at `pubsub_agent/agent.py`:

```
# Coordinator agent
root_agent = LlmAgent(
    name="pubsub_processor",
    model="gemini-2.5-flash",
    description="Event-driven document processing coordinator",
    instruction="Routes documents to specialized analyzers",
    tools=[financial_tool, technical_tool, sales_tool, marketing_tool],
)

# Sub-agents (financial, technical, sales, marketing)
# Each configured with output_schema for structured JSON
```

## Output Schemas

All sub-agents return structured Pydantic models:

```
# Financial documents return:
FinancialAnalysisOutput(
    summary: DocumentSummary,
    entities: EntityExtraction,
    financial_metrics: FinancialMetrics,
    fiscal_periods: list[str],
    recommendations: list[str]
)

# Technical documents return:
TechnicalAnalysisOutput(
    summary: DocumentSummary,
    entities: EntityExtraction,
    technologies: list[str],
    components: list[str],
    recommendations: list[str]
)

# Similar for Sales and Marketing analyzers
```



## Example Usage

### Locally without GCP:

```
cd tutorial_implementation/tutorial34
make test
```

### Test the agent in code:

```
import asyncio
from google.adk import Runner
from google.adk.sessions import InMemorySessionService
from google.genai import types
from pubsub_agent.agent import root_agent

async def test_document_analysis():
    session_service = InMemorySessionService()
    runner = Runner(
        app_name="document_analyzer",
        agent=root_agent,
        session_service=session_service
    )

    session = await session_service.create_session(
        app_name="document_analyzer",
        user_id="test_user"
    )

    prompt = types.Content(
        role="user",
        parts=[types.Part(
            text="Analyze: Revenue $1.2M, Profit 33%, Q4 2024"
        )]
    )

    async for event in runner.run_async(
        user_id="test_user",
        session_id=session.id,
        new_message=prompt
    ):
        print("Response:", event)

asyncio.run(test_document_analysis())
```

### Using ADK Web Interface:

```
adk web
```

Then visit `http://localhost:8000` and select `pubsub_processor` from the agent dropdown.

---

## Running Locally

---

### | Without Pub/Sub (Local Testing)

```
cd tutorial_implementation/tutorial34

# Run all tests
make test

# See test coverage
make test-cov
```

Tests validate:

- Agent configuration
- Sub-agent setup
- Pydantic output schemas
- Agent imports and structure

### | With Pub/Sub (Google Cloud)

After setting up GCP (see Prerequisites), run publisher and subscriber:

#### Terminal 1 - Start subscriber:

```
python subscriber.py
```

#### Terminal 2 - Publish documents:

```
python publisher.py
```

The subscriber will process each document with the coordinator agent.

---

## Google Cloud Deployment

---

### | Step 1: Set Up Pub/Sub Resources

```
gcloud pubsub topics create document-uploads
gcloud pubsub subscriptions create document-processor \
  --topic=document-uploads \
  --ack-deadline=600
```

### | Step 2: Run Subscriber

```
python subscriber.py
```

### | Step 3: Publish Documents

```
python publisher.py
```

The subscriber will automatically process each Pub/Sub message using the coordinator agent.

---

# Troubleshooting

---

## | Common Issues

### Issue 1: gcloud command not found

**Cause:** Google Cloud CLI not installed

**Solution:**

```
# macOS
brew install --cask google-cloud-sdk

# After installation, verify
gcloud --version
```

---

### Issue 2: Agent not found when running locally

**Cause:** Agent module not properly installed

**Solution:**

```
cd tutorial_implementation/tutorial34

# Install in development mode
pip install -e .

# Verify agent imports
python -c "from pubsub_agent.agent import root_agent; print(root_agent.name)"
```

---

### Issue 3: Tests fail with import errors

**Cause:** Dependencies not installed

**Solution:**

```
cd tutorial_implementation/tutorial34

# Install dependencies
make setup

# Or manually
pip install -r requirements.txt

# Run tests
make test
```

---

## Issue 4: Messages Not Delivered on Pub/Sub

**Cause:** Subscription not receiving published messages

**Solution:**

```
# Verify subscription exists
gcloud pubsub subscriptions list

# Check subscription details
gcloud pubsub subscriptions describe document-processor

# Manually pull a message to test
gcloud pubsub subscriptions pull document-processor --limit=1

# Check IAM permissions
gcloud pubsub subscriptions get-iam-policy document-processor
```

---

## Issue 5: Pub/Sub Authentication Error

**Error:** DefaultCredentialsError: Could not automatically determine credentials

**Solution:**

```
# Set up application default credentials
gcloud auth application-default login

# Or set explicit credentials

# Verify setup
gcloud auth list
```

---

## Issue 6: Tests fail with "GOOGLE\_API\_KEY not set"

**Cause:** Gemini API key not configured

**Solution:**

```
# Set your Gemini API key

# Verify it's set
echo $GOOGLE_API_KEY

# Run tests again
make test
```

---

## Issue 7: Agent processes documents but returns empty results

**Cause:** Model not returning expected output format

**Solution:**

- Verify GOOGLE\_API\_KEY is set and valid
- Check that the document content is clear and valid
- Review agent instructions in `pubsub_agent/agent.py`
- Test with a simple document first

```
# Test the agent directly
import asyncio
from google.adk import Runner
from google.adk.sessions import InMemorySessionService
from google.genai import types
from pubsub_agent.agent import root_agent

async def test():
    session_service = InMemorySessionService()
    runner = Runner(
        app_name="test",
        agent=root_agent,
        session_service=session_service
    )
    session = await session_service.create_session(
        app_name="test",
        user_id="test"
    )
    message = types.Content(
        role="user",
        parts=[types.Part(text="Revenue $1M, Profit 30%")]
    )
    async for event in runner.run_async(
        user_id="test",
        session_id=session.id,
        new_message=message
    ):
        print(event)

asyncio.run(test())
```

## Next Steps

**You've Mastered Event-Driven Agents with Pub/Sub!** 🎉

You now know how to:

- ✓ Build multi-agent coordinator systems
- ✓ Use Pydantic for structured JSON output
- ✓ Implement async agent processing
- ✓ Route documents to specialized analyzers
- ✓ Use Google Cloud Pub/Sub for event-driven processing
- ✓ Test agents locally without GCP
- ✓ Deploy to production with Pub/Sub integration

## | Key Patterns Learned

- **Coordinator + Specialist:** One agent routes to many specialized agents
- **Structured Output:** Pydantic models enforce JSON schemas
- **Async Processing:** Non-blocking document analysis
- **Event-Driven:** Pub/Sub handles message buffering and retries
- **Tool Composition:** Sub-agents as tools within coordinator

## | Continue Learning

**Tutorial 29:** UI Integration Overview

Compare all integration approaches (Next.js, Vite, Streamlit, etc.)

**Tutorial 30:** Next.js + CopilotKit Integration

Build real-time chat interfaces with React

**Tutorial 35+:** Advanced Patterns

Master deployment, scaling, and production optimization

## | Additional Resources

- [Google Cloud Pub/Sub Documentation](https://cloud.google.com/pubsub/docs) (https://cloud.google.com/pubsub/docs)
- [ADK Documentation](https://google.github.io/adk-docs/) (https://google.github.io/adk-docs/)
- [Pydantic Documentation](https://docs.pydantic.dev/) (https://docs.pydantic.dev/)
- [Gemini API Reference](https://ai.google.dev/docs) (https://ai.google.dev/docs)



**Tutorial 34 Complete!**



You've successfully built an event-driven document processing system with a multi-agent coordinator architecture. This pattern scales to millions of documents while maintaining structured, validated output.

---

**Questions or feedback?** Open an issue on the [ADK Training Repository](https://github.com/raphaelmansuy/adk_training) ([https://github.com/raphaelmansuy/adk\\_training](https://github.com/raphaelmansuy/adk_training)).

---

Generated on 2025-10-19 17:57:28 from 34\_pubsub\_adk\_integration.md

Source: Google ADK Training Hub