

TIL: Pause and Resume Invocations with Google ADK 1.16.0

Date: 2025-10-20

Difficulty: intermediate

Reading Time: 10 minutes

Tags: til, quick-learn, pause-resume, adk-1.16, state-checkpointing, fault-tolerance, human-in-loop

Description: Quick guide to using Pause and Resume Invocations: checkpoint agent state and resume execution later for long-running workflows, human-in-the-loop, and fault tolerance

TIL: Pause and Resume Invocations - Resilient Agent Workflows


| Why Pause/Resume Invocations Matter





The Problem: Long-running agent tasks need to be interrupted gracefully or paused for human feedback without losing progress. System failures can interrupt execution mid-task, causing work to be lost.

In one sentence: Pause and Resume Invocations let agents checkpoint their state at key points and resume execution later without losing context.

| Why Should You Care?

Problems it solves:

-  **Fault tolerance** - System failures don't cause work loss; resume from checkpoint

-  **Human-in-the-loop** - Agent pauses to request feedback, then continues
-  **Long-running tasks** - Complex workflows can pause at natural break points
-  **Multi-agent handoff** - State is preserved when handing off between agents
-  **State persistence** - Complete execution context is saved automatically

Perfect for:

- Data processing pipelines (batch jobs with checkpoints)
- Customer support escalations (pause for supervisor review)
- Research workflows (save progress between analysis steps)
- Approval workflows (pause for human decision-making)
- Resilient production systems (automatic recovery from failures)

Quick Example

```
from google.adk.apps import App, ResumabilityConfig
from google.adk.agents import Agent

# Create agent as usual
agent = Agent(
    name="long_task_agent",
    model="gemini-2.0-flash",
    description="Agent for long-running tasks",
    instruction="Complete tasks with checkpoints."
)

# Enable pause/resume support
app = App(
    name="resumable_app",
    root_agent=agent,
    resumability_config=ResumabilityConfig(is_resumable=True)
)

# Initial invocation - creates checkpoint on agent completion
async for event in runner.run_async(session=session, new_message=user_input):
    last_invocation_id = event.invocation_id
    # Event includes: end_of_agent=True, agent_state={...}

# [Later] Resume from checkpoint with new input
async for event in runner.run_async(
    session=session,
    new_message=new_input,
    invocation_id=last_invocation_id # Resume from here!
):
    # Execution continues from saved checkpoint
    pass
```

How It Works (3 Key Concepts)

1. State Checkpointing

When an agent completes, it emits an event with its state:

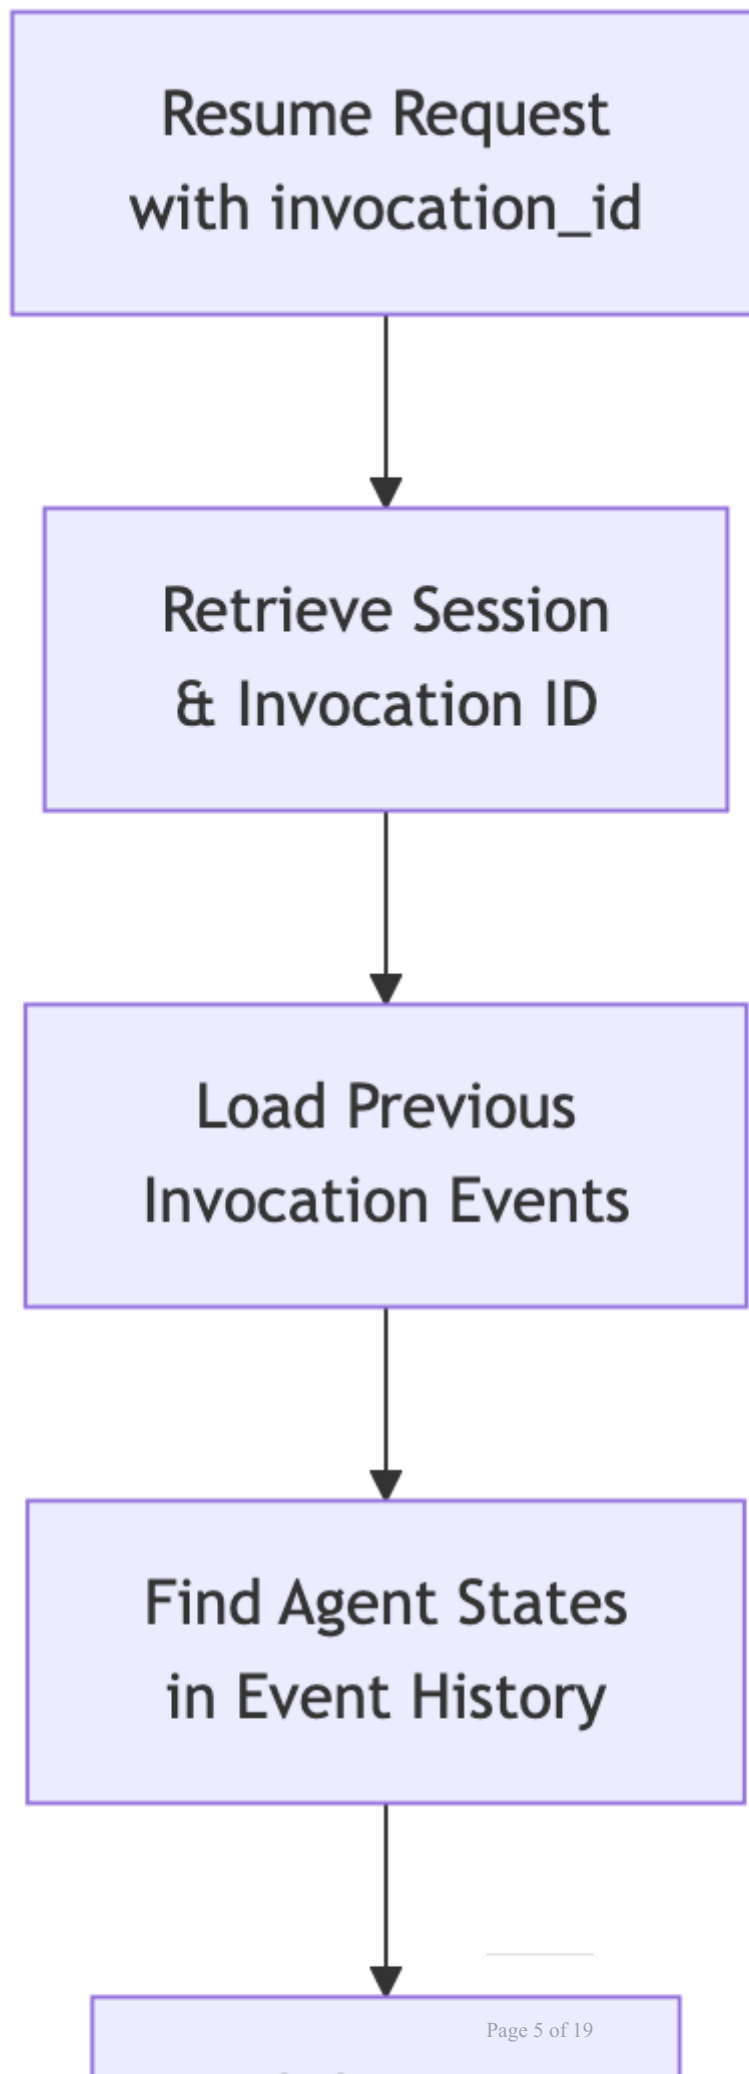
```
event = Event(  
    invocation_id='inv_1',  
    author='agent_name',  
    actions=EventActions(  
        end_of_agent=True,           # Marks completion  
        agent_state={'data': '...'} # Persisted state  
    ),  
    content=response  
)
```

The state is:

- **Automatic:** Framework handles it transparently
- **Serialized:** Converted to JSON for storage
- **Complete:** Includes all agent execution context

2. State Restoration

When resuming, the framework restores the previous state:



The restoration process:

1. **Find** previous invocation events in session
2. **Extract** agent_state from checkpoint events
3. **Restore** to InvocationContext
4. **Continue** agent execution with saved state

3. Configuration

Three simple settings to enable pause/resume:

```
from google.adk.apps import ResumabilityConfig, App

# 1. Create config
config = ResumabilityConfig(is_resumable=True)

# 2. Attach to app
app = App(root_agent=agent, resumability_config=config)

# 3. That's it! Framework handles checkpointing automatically
```

Use Cases

1. Long-Running Data Processing

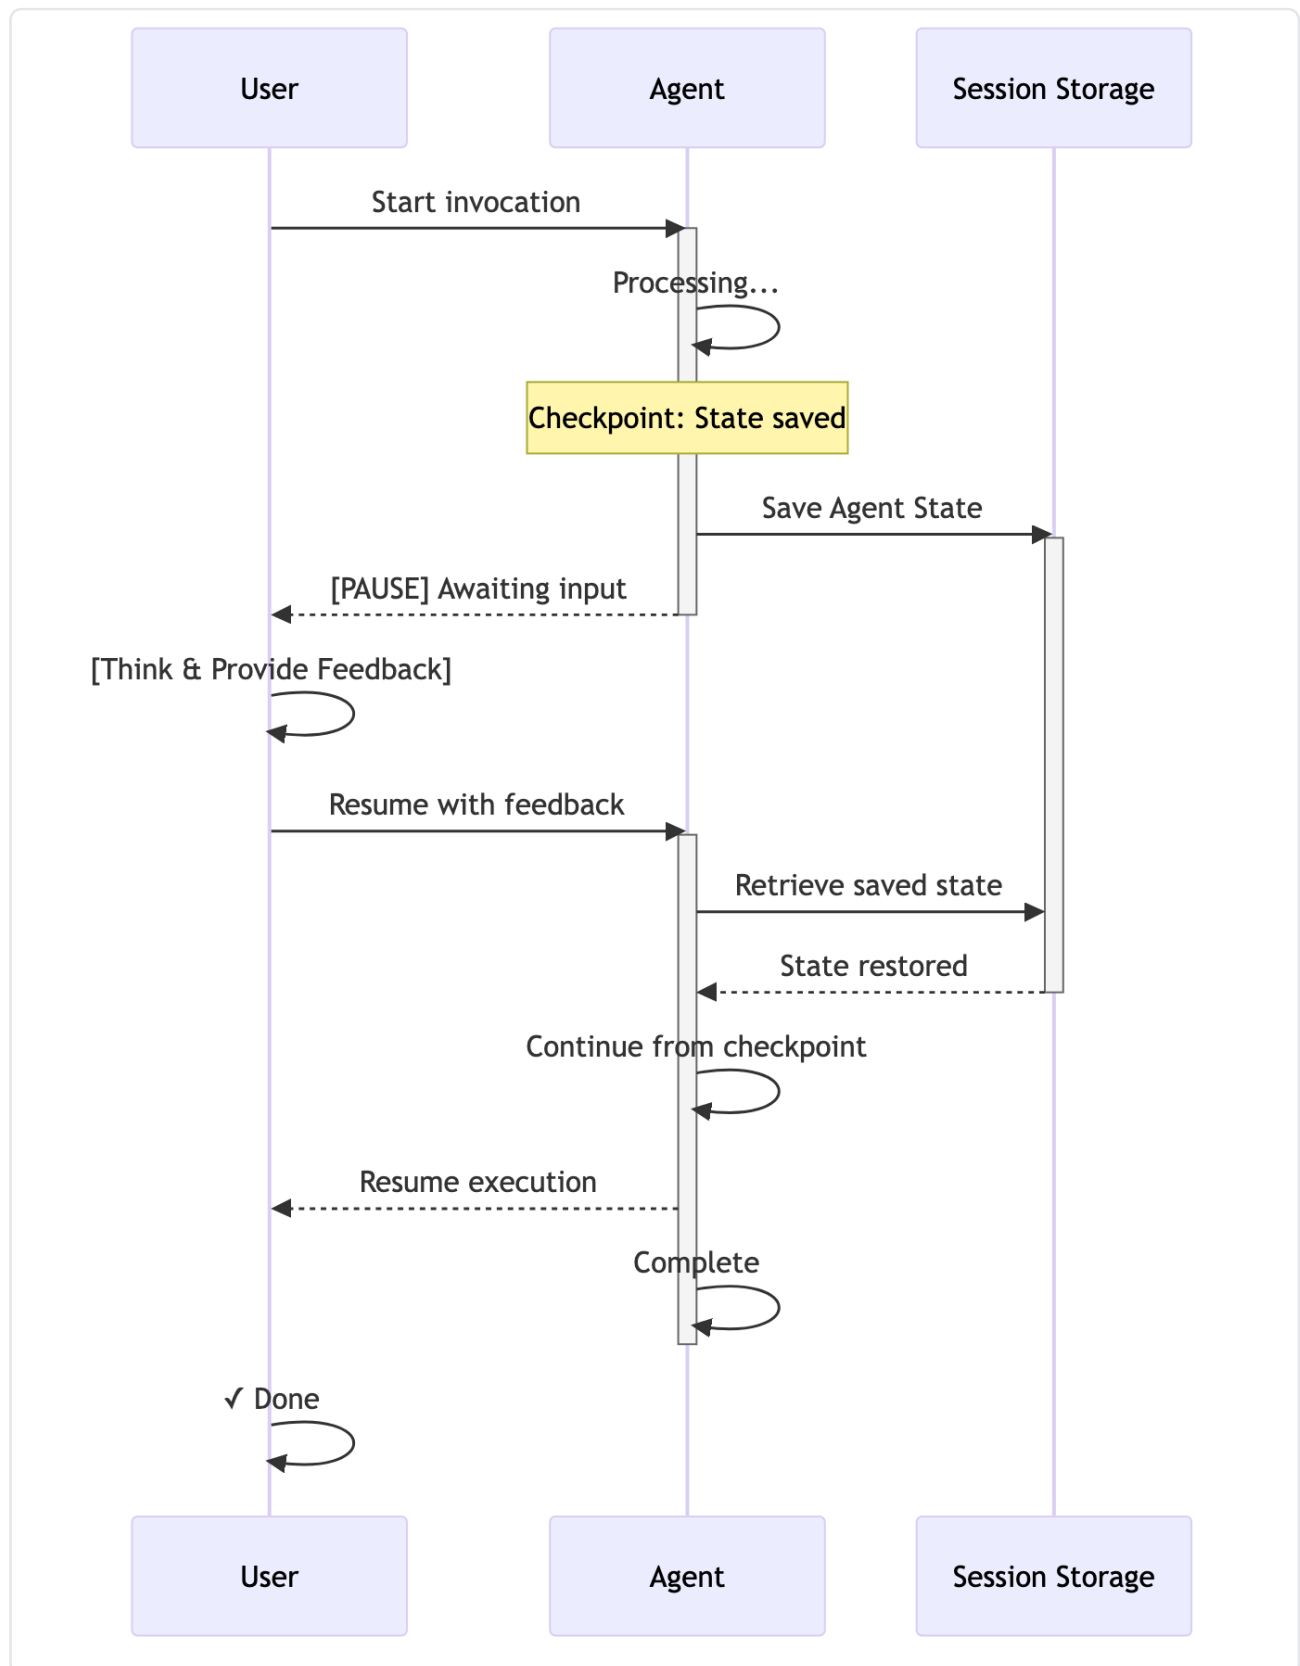
Scenario: Processing a large dataset in multiple stages

```
Stage 1: Validate Input    [CHECKPOINT 1]
      ↓
Stage 2: Process Data     [CHECKPOINT 2]
      ↓
Stage 3: Analyze Results  [CHECKPOINT 3]
      ↓
Stage 4: Generate Report  [CHECKPOINT 4]
      ↓
      Complete
```

If the system crashes after Stage 2, just resume with `invocation_id` from Stage 2's checkpoint.

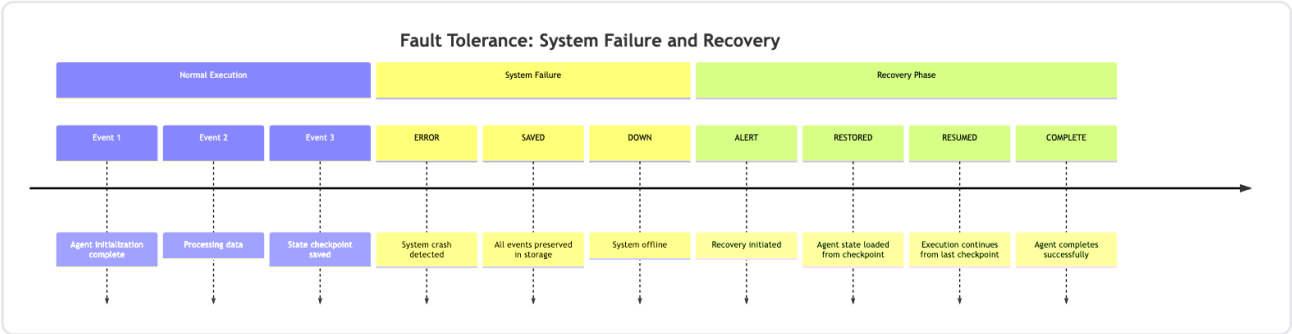
2. Human-in-the-Loop Approval

Scenario: Agent prepares decision, waits for human approval



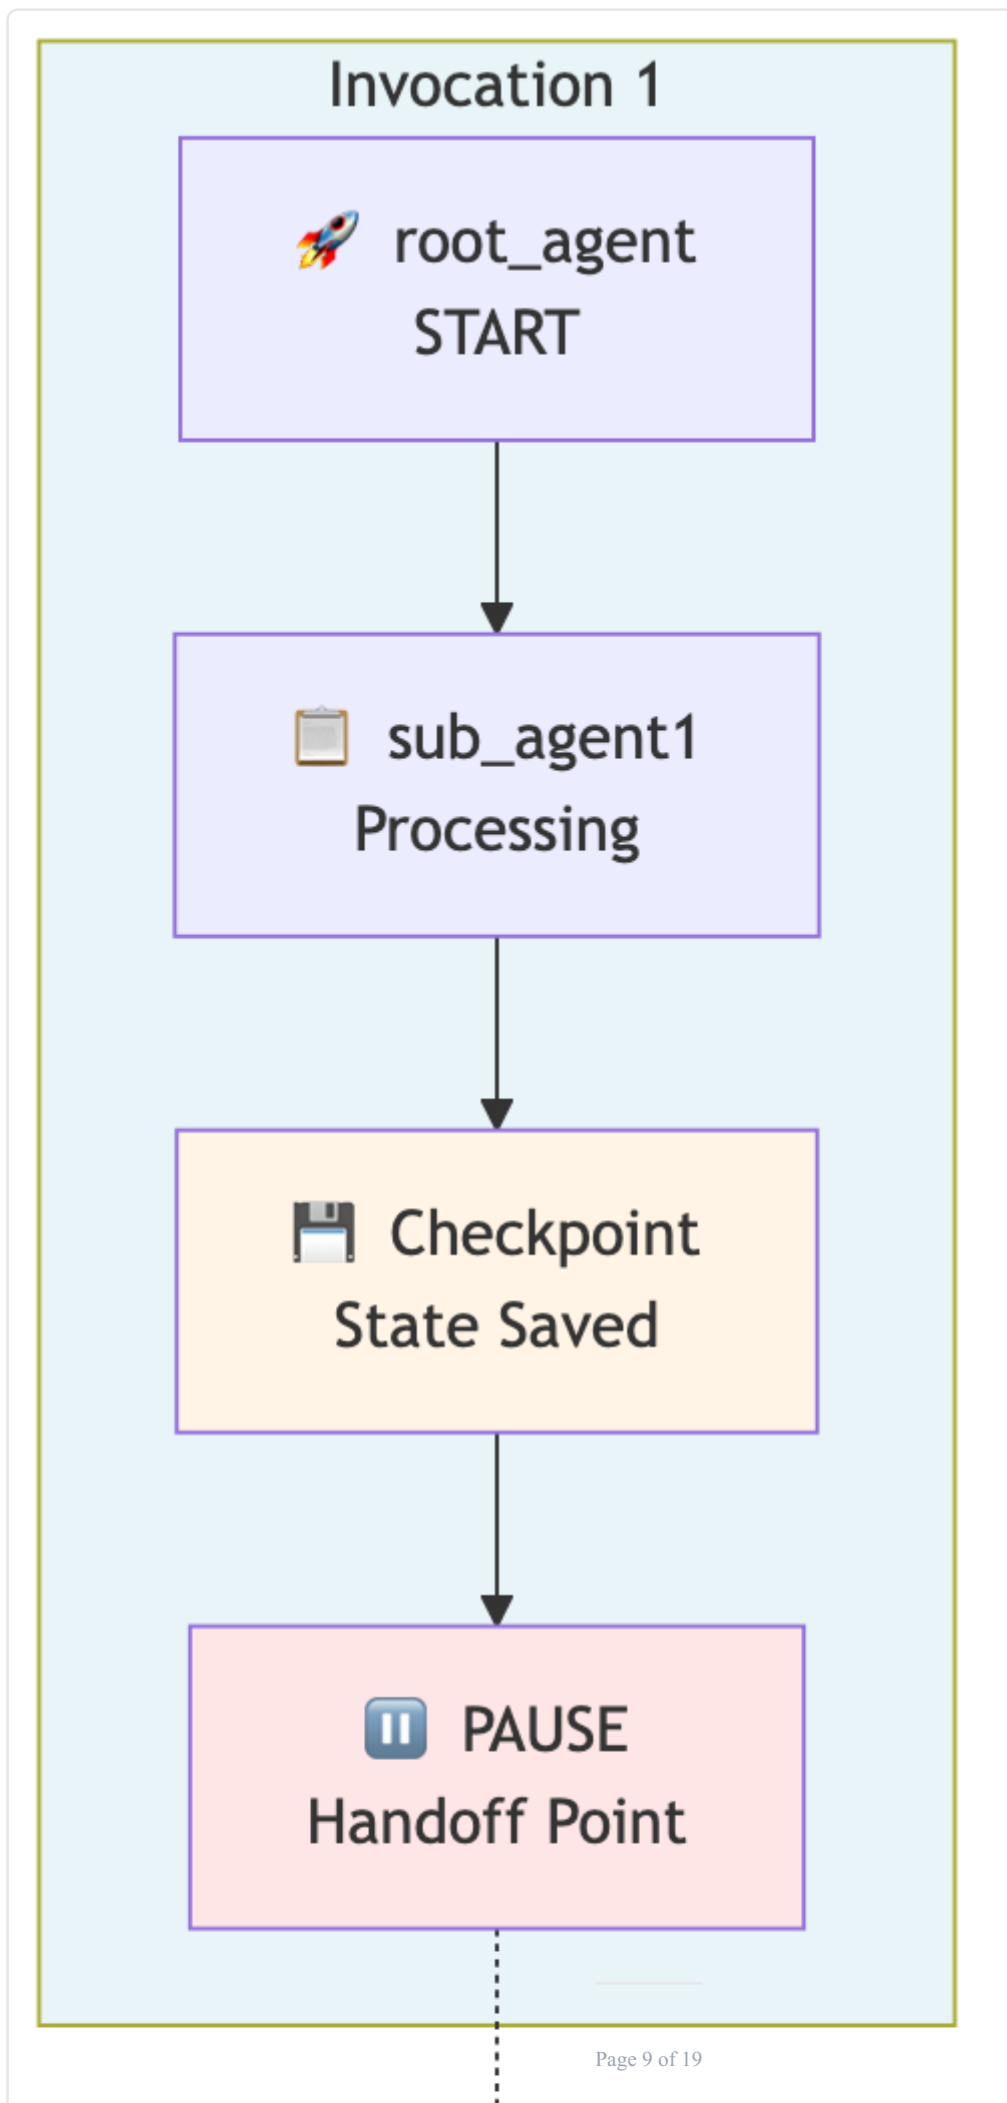
3. Fault Tolerance

Scenario: Production system with failures



4. Multi-Agent Workflows

Scenario: Sequential agent handoff with state preservation



Each agent checkpoint includes full state for potential resumption.

Key Features

ResumabilityConfig

```
config = ResumabilityConfig(  
    is_resumable=True # Enable pause/resume support  
)
```

That's all! The framework handles:

- State serialization
- Checkpoint creation
- State restoration on resume
- Event history management

Agent State Types

Different agent types have specialized states:

- **LoopAgent:** `current_sub_agent`, `times_looped`
- **SequentialAgent:** `completed_agents`, `current_index`
- **ParallelAgent:** `agent_states`, `completion_status`

Custom agents can implement their own state via:

```
class MyAgent(BaseAgent):  
    def get_current_state(self) -> dict:  
        return {  
            'progress': self.progress,  
            'data': self.accumulated_data,  
            'timestamp': datetime.now()  
        }
```

Resumption with Optional New Input

Resume with new user input:

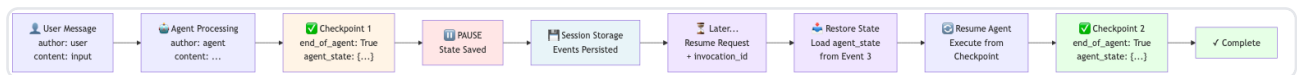
```
await runner.run_async(
    session=session,
    new_message="New feedback",          # Optional
    invocation_id=previous_invocation_id # Required
)
```

Or resume with previous input:

```
await runner.run_async(
    session=session,
    new_message=None,                # None = reuse
    invocation_id=previous_invocation_id
)
```

Event Flow Example

Events flow through the session with checkpoints marked for potential resumption:



Timeline visualization:

```

Session.events = [
  Event 1: User Message (author: 'user')
    content: { text: "Process this data" }
    agent_state: None
    |
    |
  Event 2: Agent Processing
    |   author: 'agent'
    |   content: { text: "Processing..." }
    |
    v
  Event 3: Agent Complete [CHECKPOINT]
    author: 'agent'
    actions: { end_of_agent: True }
    agent_state: { "state_key": "state_value" }
    ---> SAVED IN SESSION STORAGE <---
    |
    | [PAUSE - Can Resume Here]
    |
    v
  Event 4: Resume Point
    [Later] Resumption with invocation_id
    |
    v
  Event 5: Agent Continues
    author: 'agent'
    agent_state: RESTORED from Event 3
    |
    v
  Event 6: Agent Complete [CHECKPOINT]
    author: 'agent'
    actions: { end_of_agent: True }
    agent_state: { "agent2_state": "..." }
    ---> SAVED IN SESSION STORAGE <---
]

```

Architecture Overview

New Components:

- `ResumabilityConfig`: Configuration class
- `BaseAgentState`: Abstract state base class
- `LoopAgentState`, `SequentialAgentState`, `ParallelAgentState`: Specialized states
- `Runner._setup_context_for_resumed_invocation()`: Resumption logic

Enhanced Components:

- `App`: Now accepts `resumability_config`
- `InvocationContext`: Populates agent states from events
- `EventActions`: Includes `agent_state` field
- `Event`: Can carry agent state information

| Testing Your Implementation

The implementation includes comprehensive tests:

```
# Run all tests
pytest tests/ -v

# Run specific test
pytest tests/test_agent.py::TestAgentConfiguration -v

# With coverage
pytest tests/ --cov=pause_resume_agent
```

Expected test patterns:

```
# Test both resumable and non-resumable modes
@pytest.mark.parametrize('resumable', [True, False])
async def test_pause_resume(resumable: bool):
    config = ResumabilityConfig(is_resumable=resumable)
    app = App(root_agent=agent, resumability_config=config)
    # Test execution and checkpoint handling
```

| Best Practices

1. Always enable ResumabilityConfig if you need pause/resume

```
python config = ResumabilityConfig(is_resumable=True)
```

2. Understand your checkpoint points - Know where your agent naturally completes and saves state

3. Test resumption scenarios - Test both:

4. Normal execution from start

5. Resumption from checkpoint

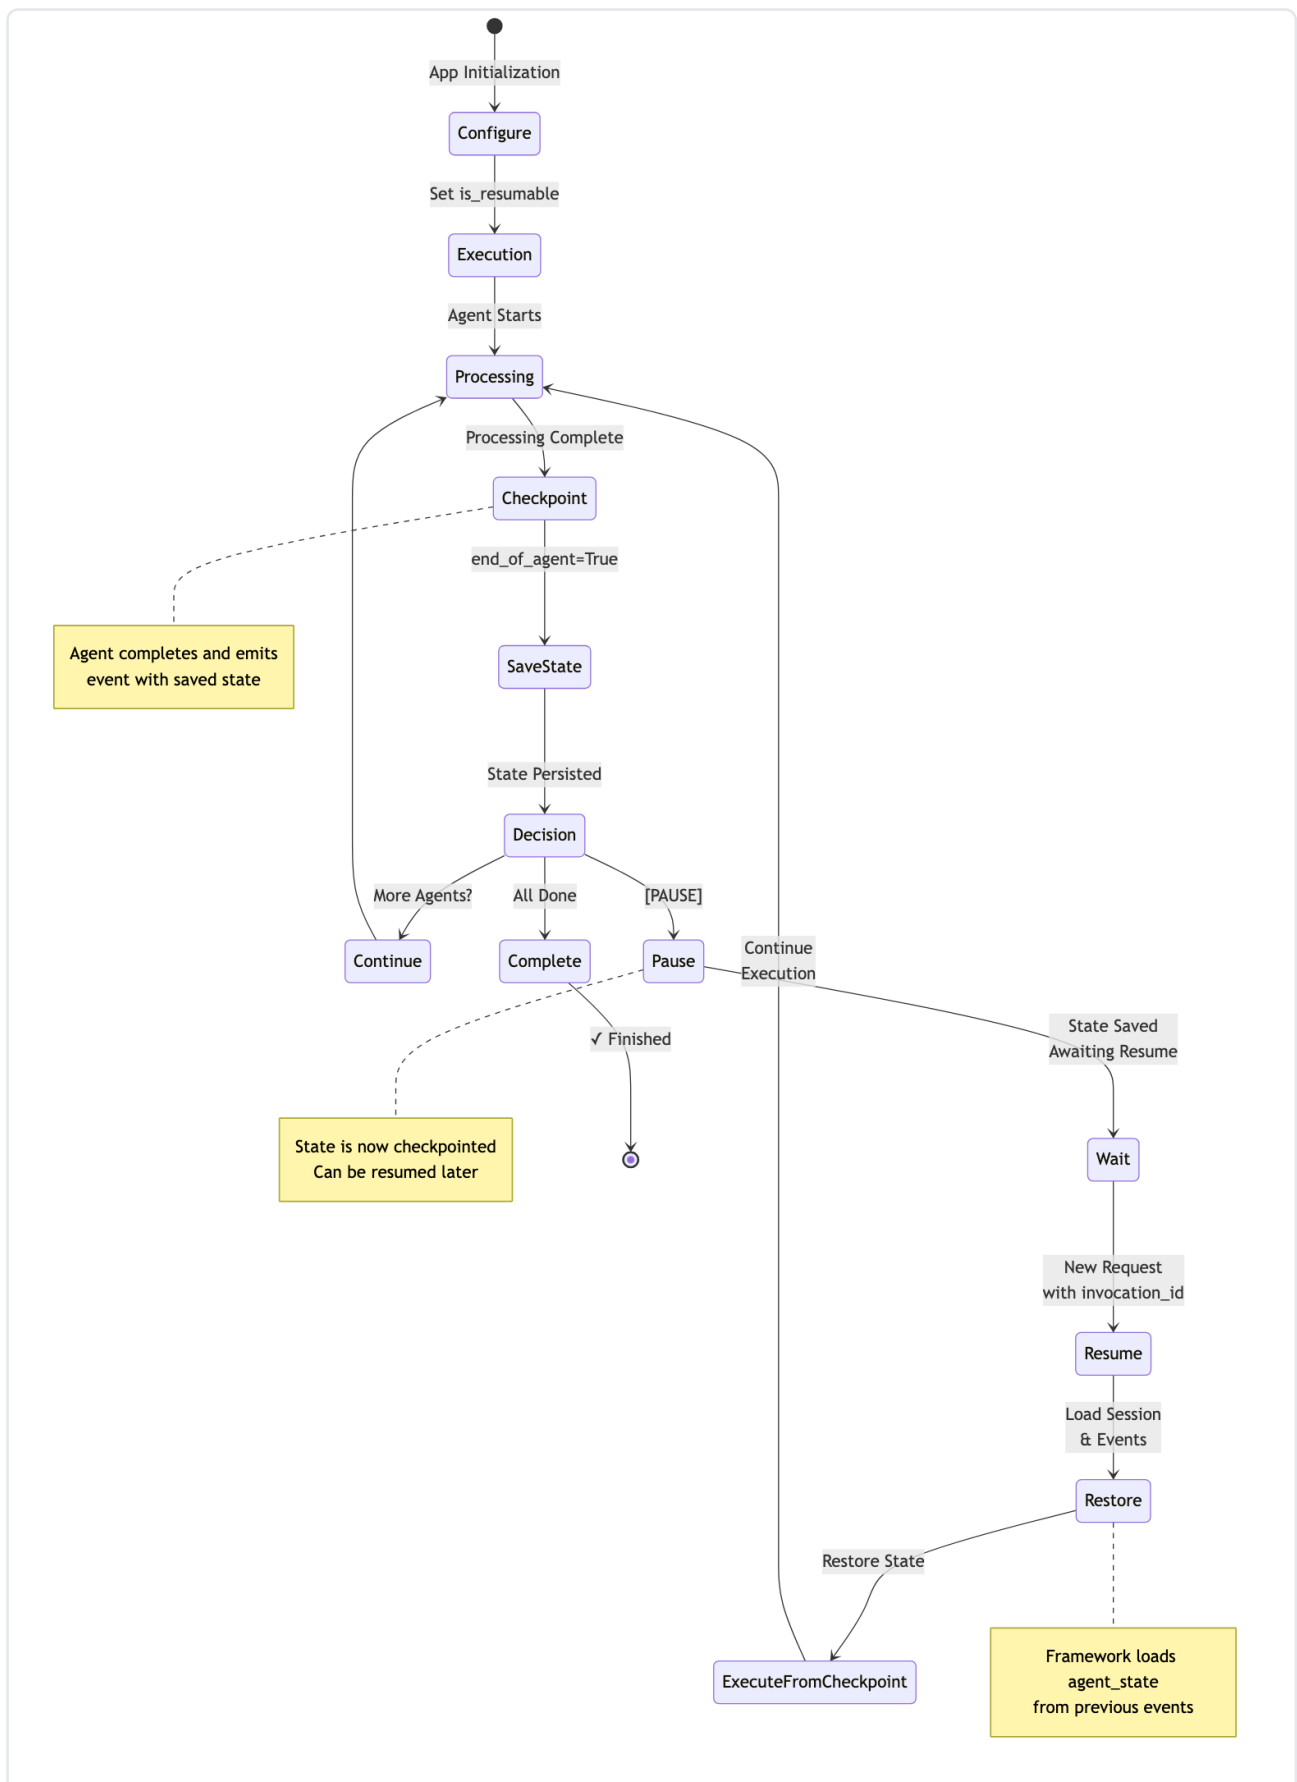
6. Handle state errors gracefully

```
python try: await runner.run_async(session, invocation_id=prev_id) except
StateRestorationError: # Fallback: start fresh await
runner.run_async(session, new_message=original_input)
```

7. Clean up old sessions - Archive sessions to avoid accumulation over time

| State Lifecycle

Understanding the state transition flow is crucial for building reliable pause/resume workflows:



Common Patterns

Pattern 1: Simple Pause/Resume

```
# Initial execution
async for event in runner.run_async(session, new_message=user_input):
    invocation_id = event.invocation_id

# Later: Resume
async for event in runner.run_async(
    session=session,
    new_message=new_input,
    invocation_id=invocation_id
):
    pass
```

Pattern 2: Human-in-the-Loop

```
# Agent pauses for feedback
await runner.run_async(session, new_message="Start process")

# Human reviews and provides input
human_feedback = request_human_input()

# Resume with feedback
await runner.run_async(
    session=session,
    new_message=human_feedback,
    invocation_id=paused_invocation_id
)
```


Pattern 3: Fault-Tolerant Processing

```
try:
    invocation_id = None
    async for event in runner.run_async(session, new_message=data):
        invocation_id = event.invocation_id
except Exception as e:
    logger.error(f"Failed, resuming from {invocation_id}")
    # Resume from checkpoint after fix
    async for event in runner.run_async(
        session=session,
        invocation_id=invocation_id
    ):
        pass
```

Limitations & Considerations

1. **App Configuration Required** - Must explicitly set `is_resumable=True`
2. **JSON Serialization** - Agent state must be JSON-serializable
3. **Session Storage** - Resuming requires session events from original invocation
4. **Sub-Agent Resumption** - Current limitation for resuming mid-sub-agent (documented in ADK)

Related Features

- **Context Caching** (v1.15.0): Complement with context preservation
- **Session Management**: VertexAiSessionService, DatabaseSessionService
- **Event Streaming**: Observe checkpoint events in real-time

Implementation Example

See the companion implementation: [til_pause_resume_20251020](https://github.com/raphaelmansuy/adk_training/tree/main/til_implementation/til_pause_resume_20251020) (https://github.com/raphaelmansuy/adk_training/tree/main/til_implementation/til_pause_resume_20251020)

Features:

- ✓ Full agent implementation with checkpoint-aware tools
- ✓ 19 comprehensive tests

- ✓ Make commands for setup, test, dev, demo
- ✓ Example tools: data processing, checkpoint validation, resumption hints
- ✓ Complete documentation in README

Quick Start with Example

```
# Clone/enter the example
cd til_implementation/til_pause_resume_20251020

# Setup
make setup

# Add API key
echo "GOOGLE_API_KEY=your_key" >> pause_resume_agent/.env

# Run tests
make test

# Launch web interface
make dev
```

References

- **ADK GitHub:** <https://github.com/google/adk-python>
- **v1.16.0 Release:** <https://github.com/google/adk-python/compare/v1.15.1...v1.16.0>
- **Related Commits:** ce9c39f, 2f1040f, 1ee01cc, f005414, fbf7576

Summary

Pause and Resume Invocations in ADK v1.16.0 enable building **resilient, interactive agent systems** by:

- ✓ Checkpointing agent state automatically at completion points
- ✓ Enabling graceful pauses for human feedback or system recovery
- ✓ Restoring full execution context on resumption

- ✓ Supporting long-running workflows without state loss
- ✓ Providing fault tolerance through state persistence

This feature is essential for production agents that need reliability, human oversight, or multi-step processing workflows.

See Also

| Related TILs

- **[TIL: Context Compaction \(/docs/til/til_context_compaction_20250119\)](/docs/til/til_context_compaction_20250119)** - Reduce token usage in long-running paused workflows
- **[Back to TIL Index \(/docs/til/til_index\)](/docs/til/til_index)** - Browse all quick-learn guides

| Related Blog Posts

- **[Deploy AI Agents: Production Strategies \(/blog/deploy-ai-agents\)](/blog/deploy-ai-agents)** - Understand resilience and fault tolerance in production systems
 - **[The Multi-Agent Pattern: Managing Complexity \(/blog/multi-agent-pattern-complexity-management\)](/blog/multi-agent-pattern-complexity-management)**
 - Understand state management across complex workflows
-
-

Generated on 2025-10-20 20:49:20 from til_pause_resume_20251020.md

Source: Google ADK Training Hub